Understanding Synchronisation

Jonathan Lawrence and Gavin Lowe

January 21, 2025

Abstract

We study *synchronisation objects*: objects that allow two or more threads to synchronise. We define a correctness condition for such synchronisation objects, which we call *synchronisation linearisation*: informally, the synchronisations appear to take place in a one-at-a-time order, consistent with the calls and returns of operations on the object. We also define a progress condition, which we call *synchronisation progressibility*: informally, executions of operations don't get stuck unnecessarily.

We consider testing of implementations of synchronisation objects. The basic idea is to run several threads that use the object, record the history of operation calls and returns, and then to test whether the resulting history satisfies synchronisation linearisation and progressibility. We present algorithms for this last step, and also present results concerning the complexity of the problem.

1   Introduction

In many concurrent programs, it is necessary at some point for two or more threads to *synchronise*: each thread waits until other relevant threads have reached the synchronisation point before continuing; in addition, the threads can exchange data. Reasoning about programs can be easier when synchronisations are used: it helps to keep threads in consistent stages of the program, and so makes it easier to reason about the states of different threads.

We study synchronisations in this paper: we describe how synchronisations can be specified, and what it means for such a specification to be satisfied. We also describe techniques for testing implementations.

We start by giving some examples of synchronisations in order to illustrate the idea. (We use Scala notation; we explain non-standard aspects of the

language in footnotes.) In each case, the synchronisation is mediated by a *synchronisation object*.

Perhaps the most common form of synchronisation object is a synchronous channel. Such a channel might have signature[1]

```scala
class SyncChan[A]{
  def send(x: A): Unit
  def receive(): A
}
```

Each execution of one of the operations must synchronise with an execution of the other operation: the two executions must overlap in time. If an execution send(x) synchronises with an execution of receive, then the receive returns x.

Each synchronisation of a synchronous channel involves two executions of *different* operations; we say that the synchronisation is *heterogeneous*. By contrast, sometimes two executions of the *same* operation may synchronise; we say that the synchronisation is *homogeneous*. For example, an *exchanger* has the following signature.

```scala
class Exchanger[A]{
  def exchange(x: A): A
}
```

When two threads call exchange, the executions can synchronise, and each receives the value passed in by the other.

For some synchronisation objects, synchronisations might involve more than two threads. For example, a *barrier synchronisation* object of the following class

```scala
class Barrier(n: Int){
  def sync(me: Int): Unit
}
```

can be used to synchronise n threads. Each thread is assumed to have an integer thread identifier in the range [0..n). Each thread me calls sync(me), and no execution returns until all n have called it. We say that the synchronisation has *arity* n.

A *combining barrier*, in addition to acting as a barrier synchronisation, also allows each thread to submit a parameter, and for all to receive back some function of those parameters.[2]

```scala
class CombiningBarrier[A](n: Int, f: (A,A) => A){
```

---

[1] The class is polymorphic in the type A of data. The type Unit is the type that contains a single value, the *unit value*, denoted ().

[2] The Scala type (A,A) => A represents functions from pairs of A to A.

```
  def sync(x: A): A
}
```

The function f is assumed to be associative. If n threads call sync
with parameters $x_1, \ldots, x_n$, in some order, then each receives back
$f(x_1, f(x_2, \ldots f(x_{n-1}, x_n) \ldots))$. (In the common case that f is commutative,
this result is independent of the order of the parameters.)

Some synchronisation objects allow different modes of synchronisation.
For example, consider a synchronous channel with timeouts: each execu-
tion might synchronise with another execution, or might timeout without
synchronisation. Such a channel has a signature as follows.

```
class TimeoutChannel[A]{
  def send(x: A): Boolean
  def receive(): Option[A]
}
```

The send operation returns a boolean to indicate whether the send was suc-
cessful, i.e. whether it synchronised. The receive operation can return a value
Some(x) to indicate that it synchronised and received x, or can return the
value None to indicate that it failed to synchronise[3]. Thus an execution of
each operation may or may not synchronise with an execution of the other
operation. Unsuccessful executions of send and receive can be considered
unary synchronisations.

So far, our example synchronisation objects have been *stateless*: they
maintain no state from one synchronisation to another. By contrast, some
synchronisation objects are *stateful*: they maintain some state between syn-
chronisations, which might affect the availability of the results of synchroni-
sations. As a toy example, consider a synchronous channel that, in addition,
maintains a sequence counter, and such that both executions receive the
value of this counter.

```
class SyncChanCounter[A]{
  private var counter: Int
  def send(x: A): Int      // result is sequence number
  def receive(): (A, Int)  // result is (value received, sequence number)
}
```

Some implementations of synchronous channels allow the channel to be
closed, say by a unary operation close.

```
class CloseableChan[A]{
  def send(x: A): Unit
  def receive(): A
```

---

[3]The type Option[A] contains the union of such values.

3

```
  def close(): Unit
}
```

Calls to send or receive after the channel is closed throw an exception. Thus such an object is stateful, with two states, open and closed; and the operations have different modes of synchronisation, either successful or throwing an exception.

An *enrollable barrier* (based on [WBM+10]) is a barrier that allows threads to enrol and resign (via unary operations):

```
class EnrollableBarrier(n: Int){
  def sync(me: Int): Unit
  def enrol(me: Int): Unit
  def resign(me: Int): Unit
}
```

Each barrier synchronisation is between all threads that are currently enrolled, so sync has a variable arity. The barrier has a state, namely the currently enrolled threads.

A *terminating queue* can also be thought of as a stateful synchronisation object with multiple modes. Such an object mostly acts like a standard partial concurrent queue: if a thread attempts to dequeue, but the queue is empty, it blocks until the queue becomes non-empty. However, if a state is reached where all the threads are blocked in this way, then they all return a special value to indicate this fact. In some concurrent algorithms, such as a concurrent graph search, this latter outcome indicates that the algorithm should terminate. Such a terminating queue might have the following signature, where a dequeue returns the value None to indicate the termination case.

```
class TerminatingQueue[A](n: Int){ // n is the number of threads
  def enqueue(x: A): Unit
  def dequeue: Option[A]
}
```

The termination outcome can be seen as a synchronisation between all n threads. This terminating queue combines the functionality of a concurrent datatype and a synchronisation object.

In this paper, we consider what it means for one of these synchronisation objects to be correct. We also present techniques for testing correctness of implementations.

In Section 2 we describe how to specify a synchronisation object. The definition has similarities with the standard definition of *linearisation* for concurrent datatypes, except it talks about synchronisations between executions

4

of operations, rather than single executions: we call the property *synchronisation linearisation*; informally, the synchronisations appear to take place in a one-at-a-time order, consistent with the calls and returns of operations on the object. We also define a progress condition, which we call *synchronisation progressibility*: informally, executions don't get stuck unnecessarily.

In Section 3 we consider the relationship between synchronisation linearisation and (standard) linearisation. We show that linearisation is an instance of synchronisation linearisation, but that synchronisation linearisation is more general. We also show that synchronisation linearisation corresponds to a small adaptation of linearisation, where an operation of the synchronisation object may correspond to *two* operations of the object used to specify linearisation.

We then consider testing of synchronisation object implementations. Our techniques are based on the techniques for testing (standard) linearisation [Low16], which we sketch in Section 4: the basic idea is to record a history of threads using the object, and then to test whether that history is linearisable. In Section 5 we show how the technique can be adapted to test for synchronisation linearisation, using the result of Section 3. Then in Section 6 we show how synchronisation linearisation can be tested more directly: we describe algorithms that test whether a history of a synchronisation object is synchronisation linearisable. We also present various complexity results: testing whether a history is synchronisation linearisable is NP-complete in general, but can be solved in polynomial time in the case of binary (heterogeneous or homogeneous) stateless synchronisation objects.

In Section 7 we describe experiments to determine the effectiveness of the testing techniques. We sum up in Section 8.

$\boxed{\text{Contributions}}$

2 Specifying synchronisations

In this section we describe how synchronisations can be formally specified. We start by considering *heterogeneous binary* synchronisation, i.e. where every synchronisation is between *two* executions of *different* operations. We allow stateful synchronisation objects (which includes stateless objects as degenerative cases). We generalise later in this section.

We assume that the synchronisation object has two operations, each of which has a single parameter, with signatures as follows.

def $op_1(x_1\colon A_1)\colon B_1$
def $op_2(x_2\colon A_2)\colon B_2$

(We can model a concrete operation that takes $k \neq 1$ parameters by an operation that takes a k-tuple as its parameter. We identify a 0-tuple with the unit value, but will sometimes omit that value in examples.) In addition, the synchronisation object might have some state. Each execution of op_1 must synchronise with an execution of op_2, and vice versa. The result of each execution may depend on the two parameters, x_1 and x_2, and the current state. In addition, the state may be updated. The external behaviour is consistent with the synchronisation happening atomically at some point within the duration of both operation executions (which implies that the executions must overlap): we refer to this point as the *synchronisation point*.

Synchronisation linearisation is defined in terms of a *synchronisation specification object*: we define these specification objects in the next subsection. In Section 2.2, we review the notion of linearisation, on which synchronisation linearisation is based. We then define synchronisation linearisation, for binary heterogeneous synchronisation objects, in Section 2.3. We generalise to other classes of synchronisation objects in Section 2.4. We present our progress property, synchronisation progressibility, in Section 2.5.

2.1 Synchronisation specification objects

Each synchronisation object, with a signature as above, can be specified using a *synchronisation specification object* with the following signature.

class Spec{
  def sync(x_1: A_1, x_2: A_2): (B_1, B_2)
}

The idea is that if two executions $op_1(x_1)$ and $op_2(x_2)$ synchronise, then the results y_1 and y_2 of the executions are such that $sync(x_1, x_2)$ could return the pair (y_1, y_2). The specification object might have private state, which can be accessed and updated within sync. Note that executions of sync occur *sequentially*. We assume that sync is a deterministic function of its parameters and the state.

We formalise below what it means for a synchronisation object to satisfy the requirements of a synchronisation specification object. But first, we give some examples to illustrate the style of specification.

A generic definition of a specification object might take the following form:

class Spec{
  private var state: S
  def sync(x_1: A_1, x_2: A_2): (B_1, B_2) = {
    require(guard(x_1, x_2, state))

```
    val res₁ = f₁(x₁, x₂, state); val res₂ = f₂(x₁, x₂, state)
    state = update(x₁, x₂, state)
    (res₁, res₂)
  }
}
```

The object has some local state, which persists between executions. The
require clause of sync specifies a precondition for the synchronisation to take
place. The values res_1 and res_2 represent the results that should be returned
by the corresponding executions of op_1 and op_2, respectively. The function
update describes how the local state should be updated. We assume the
specification object is deterministic: f_1, f_2 and update are functions.

For example, consider a synchronous channel with operations

def send(x: A): Unit
def receive(u: Unit): A

(Note that we model the receive operation as taking a parameter of type
Unit, in order to fit our uniform setting.) This can be specified using a
synchronisation specification object with empty state:

class SyncChanSpec[A]{
  def sync(x: A, u: Unit): (Unit, A) = ((), x)
}

If send(x) synchronises with receive(()), then the former receives the unit
value (), and the latter receives x.

As another example, consider a filtering channel.

class FilterChan[A]{
  def send(x: A): Unit
  def receive(p: A => Boolean): A
}

Here the receive operation is passed a predicate p describing a required prop-
erty of any value received. This can be specified using a stateless specification
object with operation

def sync(x: A, p: A => Boolean): (Unit, A) = { require(p(x)); ((), x) }

The require clause specifies that executions send(x) and receive(p) can synchro-
nise only if p(x).

As an example illustrating the use of state in the synchronisation object,
recall the synchronous channel with a sequence counter, SyncChanCounter,
from the introduction. This can be specified using the following specification
object.

```
class SyncChanCounterSpec[A]{
  private var counter = 0
  def sync(x: A, u: Unit): (Int, (A, Int)) = {
    counter += 1; (counter, (x, counter))
  }
}
```

Each synchronisation increments the counter, and the value is returned to each thread.

2.2 Linearisability

We formalise the allowable behaviours captured by a particular synchronisation specification object. Our definition has much in common with the well known notion of *linearisation* [HW90], used for specifying concurrent datatypes; so we start by reviewing that notion. There are a number of equivalent ways of defining linearisation: we choose a way that will be convenient subsequently.

A *concurrent history* of an object o (either a concurrent datatype or a synchronisation object) records the calls and returns of operations on o. It is a sequence of events of the following forms:

- call.$op^i(x)$, representing a call of operation op with parameter x;

- return.op^i:y, representing a return of an execution of op, giving result y.

Here i is a *execution identity*, used to identify a particular execution, and to link the call and corresponding return. In order to be well formed, each execution identity must appear on at most one call event and at most one return event; and for each event return.op^i:y, the history must contain an earlier event call.$op^i(x)$, i.e. for the same operation and execution identity. We consider only well formed histories from now on.

We say that a call event and a return event *match* if they have the same execution identifier. A concurrent history is *complete* if for every call event, there is a matching return event, i.e. no execution is still pending at the end of the history.

For example, consider the following complete concurrent history of a concurrent object that is intended to implement a queue, with operations enq and deq.

$$h = \langle \text{call.enq}^1(5), \text{call.enq}^2(4), \text{call.deq}^3(), \\ \text{return.enq}^1:(), \text{return.deq}^3:4, \text{return.enq}^2:() \rangle.$$

This history is illustrated by the timeline in Figure 1.

Figure 1: Timeline representing the linearisation example. Time runs from left to right; each horizontal line represents an operation execution, with the left-hand end representing the call event, and the right-hand end representing the return event

Linearisability is specified with respect to a specification object $Spec$, with the same operations (and signatures) as the concurrent object in question. A history of the specification object is a sequence of events of the form:

- $op^i(x){:}y$ representing an execution of operation op with parameter x, returning result y; again i is an execution identity, which must appear at most once in the history.

A history is *legal* if it is consistent with the definition of $Spec$, i.e. for each execution, the precondition is satisfied, and the return value is as for the definition of the operation in $Spec$. We assume that the specification object is deterministic: after a particular history, there is a unique value that can be returned by each execution.

For example, consider the history

$$h_s \quad = \quad \langle \mathsf{enq}^2(4){:}(), \ \mathsf{enq}^1(5){:}(), \ \mathsf{deq}^3(){:}4\rangle.$$

This is a legal history for a specification object that represents a queue. This history is illustrated by the "×"s in Figure 1.

Let h be a complete concurrent history, and let h_s be a legal history of the specification object. We say that h and h_s *correspond* if they contain the same executions, i.e., for each call.$op^i(x)$ and return.$op^i{:}y$ in h, h_s contains $op^i(x){:}y$, and vice versa. We say that h_s is a *linearisation* of h if there is some way of interleaving the two histories (i.e. creating a history containing the events of h and h_s, preserving the order of events) such that each $op^i(x){:}y$ occurs between call.$op^i(x)$ and return.$op^i{:}y$. Informally, this indicates that the executions of h appeared to take place in the order described by h_s, and that this order is consistent with the specification object.

Figure 2: Timeline representing the synchronisation example.

Continuing the running example, h_s is a linearisation of h, as evidenced by the interleaving

$$\langle \mathsf{call.enq}^1(5),\ \mathsf{call.enq}^2(4),\ \mathsf{enq}^2(4){:}(),\ \mathsf{enq}^1(5){:}(),\ \mathsf{call.deq}^3(),$$
$$\mathsf{return.enq}^1{:}(),\ \mathsf{deq}^3{:}4,\ \mathsf{return.deq}^3{:}4,\ \mathsf{return.enq}^2{:}()\rangle,$$

as illustrated in Figure 1.

A concurrent history might not be complete, i.e. it might have some pending executions that have been called but have not returned. An *extension* of a history h is formed by adding zero or more return events corresponding to pending executions. We write $complete(h)$ for the subsequence of h formed by removing all call events corresponding to pending executions. We say that a (not necessarily complete) concurrent history h is *linearisable* if there is an extension h' of h such that $complete(h')$ is linearisable. We say that a concurrent object is linearisable if all of its histories are linearisable.

2.3   Synchronisation linearisability

We now adapt the definition of linearisability to synchronisations. For the moment, we consider only binary synchronisations. We consider a concurrent history of the synchronisation object $Sync$. The history contains call and return events, as in the previous subsection; in the case of binary synchronisations, the events correspond to the operations op_1 and op_2.

For example, the following is a complete history of the synchronous channel from earlier, and is illustrated in Figure 2:

$$\begin{aligned}
h \ = \ &\langle \mathsf{call.send}^1(8),\ \mathsf{call.send}^2(8),\ \mathsf{call.receive}^3(()),\ \mathsf{return.receive}^3{:}8,\\
&\mathsf{call.receive}^4(()),\ \mathsf{return.send}^1{:}(),\ \mathsf{call.send}^5(9),\ \mathsf{return.receive}^4{:}9,\\
&\mathsf{call.receive}^6(()),\ \mathsf{return.send}^2{:}(),\ \mathsf{return.send}^5{:}(),\ \mathsf{return.receive}^6{:}8\rangle.
\end{aligned}$$

A history of a synchronisation specification object $Spec$ is a sequence of events of the form $\mathsf{sync}^{i_1,i_2}(x_1,x_2){:}(y_1,y_2)$, representing an execution of sync

with parameters (x_1, x_2) and result (y_1, y_2). The event's identity is (i_1, i_2): each of i_1 and i_2 must appear at most once in the history. Informally, an event $\mathsf{sync}^{i_1, i_2}(x_1, x_2){:}(y_1, y_2)$ corresponds to a synchronisation between executions $\mathsf{op}_1^{i_1}(x_1){:}y_1$ and $\mathsf{op}_2^{i_2}(x_2){:}y_2$ in a history of the corresponding synchronisation object.

A history is *legal* if is consistent with the definition of the specification object. For example, the following is a legal history of SyncChanSpec.

$$h_s \;=\; \langle \mathsf{sync}^{1,3}(8, ()){:}((), 8),\; \mathsf{sync}^{5,4}(9, ()){:}((), 9),\; \mathsf{sync}^{2,6}(8, ()){:}((), 8) \rangle.$$

The history is illustrated by the "×"s in Figure 2: each event corresponds to the synchronisation of two operations, so is depicted by two "×"s on the corresponding operations, linked by a dashed vertical line. This particular synchronisation specification object is stateless, so in fact any permutation of this history would also be legal (but not all such permutations will be compatible with the history of the synchronisation object); but the same will not be true in general of a specification object with state.

Let h be a complete history of the synchronisation object $Sync$. Informally, we say that a legal history h_s of $Spec$ *corresponds* to h if their identities agree; more precisely:

- For each sync event with identity (i_1, i_2) in h_s, h contains an execution of op_1 with identity i_1 and an execution of op_2 with identity i_2;

- For each execution of op_1 with identity i_1 in h, h_s contains a sync event with identity (i_1, i_2) for some i_2;

- For each execution of op_2 with identity i_2 in h, h_s contains a sync event with identity (i_1, i_2) for some i_1.

Given a complete history h of $Sync$ and a corresponding legal history h_s of $Spec$, we say that h_s is a *synchronisation linearisation* of h if there is some way of interleaving h and h_s such that each event $\mathsf{sync}^{i_1, i_2}(x_1, x_2){:}(y_1, y_2)$ occurs between $\mathsf{call.op}_1^{i_1}(x_1)$ and $\mathsf{return.op}_1^{i_1}{:}y_1$, and between $\mathsf{call.op}_2^{i_2}(x_2)$ and $\mathsf{return.op}_2^{i_2}{:}y_2$. In the running example, h_s is a synchronisation linearisation of h, as shown by the interleaving in Figure 2.

Given a (not necessarily complete) concurrent history h and a corresponding legal history h_s of $Spec$, we say that h_s is a *synchronisation linearisation* of h if there is an extension h' of h such that h_s is a synchronisation linearisation of $complete(h')$. We say that h is synchronisation-linearisable in this case. We say that a synchronisation object is synchronisation-linearisable if all of its histories are synchronisation-linearisable.

11

2.4 Variations

Above we considered heterogeneous binary synchronisations, i.e. synchronisations between *two* executions of *different* operations, with a single mode of synchronisation.

It is straightforward to generalise to synchronisations between an arbitrary number of operation executions, some of which might be executions of the same operation. Consider a k-way synchronisation between operations

def $\mathsf{op}_j(\mathsf{x}_j\colon \mathsf{A}_j)\colon \mathsf{B}_j$ for $j = 1, \ldots, k,$

where the op_j might not be distinct. The specification object will have a corresponding operation

def $\mathsf{sync}(\mathsf{x}_1\colon \mathsf{A}_1,\ \ldots,\ \mathsf{x}_k\colon \mathsf{A}_k)\colon (\mathsf{B}_1,\ \ldots,\ \mathsf{B}_k)$

For example, for the combining barrier CombiningBarrier(n, f) of the Introduction, the corresponding specification object would be

```
class CombiningBarrierSpec{
  def sync(x₁: A, ..., xₙ: A) = {
    val result = f(x₁, f(x₂ ,... f(xₙ₋₁, xₙ )...));  (result ,..., result)
  }
}
```

A history of the specification object will have corresponding events $\mathsf{sync}^{i_1,\ldots,i_k}(x_1, \ldots, x_k)\colon(y_1, \ldots, y_k)$. The definition of synchronisation linearisation is an obvious adaptation of earlier: in the interleaving of the complete history of the synchronisation history and the history of the specification object, each $\mathsf{sync}^{i_1,\ldots,i_k}(x_1, \ldots, x_k)\colon(y_1, \ldots, y_k)$ occurs between $\mathsf{call.op}_1^{i_j}(x_j)$ and $\mathsf{return.op}_j^{i_j}\colon y_j$ for each $j = 1, \ldots, k$.

Note that if several of the op_j are the same operation, there is a choice as to the order in which their parameters are passed to sync. (However, in the case of the combining barrier, if f is associative and commutative, the order makes no difference.)

It is also straightforward to adapt the definitions to deal with multiple modes of synchronisation: the specification object has a different operation for each mode. For example, recall the TimeoutChannel from the Introduction, where sends and receives may timeout and return without synchronisation. The corresponding specification object would be:

```
class TimeoutChannelSpec[A]{
  def syncₛ(x: A) = false      // send times out
  def syncᵣ(u: Unit) = None  // receive times out
  def syncₛ,ᵣ(x: A, u: Unit) = (true, Some(x))  // synchronisation
}
```

The operation sync_s corresponds to a send returning without synchronising; likewise sync_r corresponds to a receive returning without synchronising; and $\mathsf{sync}_{s,r}$ corresponds to a send and receive synchronising. The formal definition of synchronisation linearisation is the obvious adaptation of the earlier definition: in particular sync_s must occur between the call and return of a send that doesn't synchronise, and likewise for sync_r.

As another example, the following is a specification object for a channel with a close operation.

```scala
class ClosableChannelSpec[A]{
  private var isClosed = false  // is the channel closed?
  def close(u: Unit) = { isClosed = true; () }
  def sync(x: A, u: Unit) = { require(!isClosed); ((), x) }
  def sendFail(x: A) = { require(isClosed); throw new Closed }
  def receiveFail(u: Unit) = { require(isClosed); throw new Closed }
}
```

A send and receive can synchronise corresponding to sync, but only before the channel is closed; or each may fail once the channel is closed, corresponding to sendFail and receiveFail.

2.5   Specifying progress

We now consider a progress condition for synchronisation objects.

We assume that each pending operation execution is scheduled infinitely often, unless it is blocked (for example, trying to obtain a lock); in other words, the scheduler is fair to each execution. Under this assumption, our progress condition can be stated informally as:

- If some pending operation executions can synchronise (according to the synchronisation specification object), then some such synchronisation should happen;

- Once a particular execution has synchronised, it should eventually return.

Note that there might be several different synchronisations possible. For example, consider a synchronous channel, and suppose there are pending calls to send(4), send(5) and receive. Then the receive could synchronise with *either* send, nondeterministically; subsequently, the receive should return the appropriate value, and the corresponding send should also return. In such cases, our progress condition allows *either* synchronisation to occur.

Our progress condition allows all pending executions to block if no synchronisation is possible. For example, if every pending executions on a synchronous channel is a send, then clearly none can return.

Note that our progress condition is somewhat different from the condition of *lock freedom* for concurrent datatypes [HS12]. That condition requires that, assuming operation executions collectively are scheduled infinitely often, then eventually some execution returns. Lock freedom makes no assumption about scheduling being fair. For example, if a particular operation execution holds a lock then lock freedom allows the scheduler to never schedule that execution; in most cases, this will mean that no execution returns: any implementation that uses a lock in a non-trivial way is not lock-free.

By contrast, our assumption, that each unblocked pending execution is scheduled infinitely often, reflects that modern schedulers *are* fair, and do not starve any single execution. For example, if an execution holds a lock, and is not in a potentially unbounded loop or permanently blocked trying to obtain a second lock, then it will be scheduled sufficiently often, and so will eventually release the lock. Thus our progress condition can be satisfied by an implementation that uses locks. However, our assumption does allow executions to be scheduled in an unfortunate order (as long as each is scheduled infinitely often), which may cause the progress condition to fail. It also allows other synchronisation primitives, such as locks and semaphores to be unfair: for example, a thread that is repeatedly trying to obtain a lock may repeatedly fail as other threads obtain the lock.

The following definition identifies maximal sequences of synchronisations that could occur given a particular history of a synchronisation object.

Definition 1 Given a history h of the synchronisation object and a legal history h_s of the specification object, we say that h_s is a *maximal synchronisation linearisation* of h if:

- h_s is a synchronisation linearisation of h;

- no proper legal extension $h_s \frown \langle e \rangle$ of h_s is a synchronisation linearisation of h.

For example, the history

$$h = \langle \mathsf{call.send}^1(4), \mathsf{call.send}^2(5), \mathsf{call.receive}^3(()) \rangle$$

has two maximal synchronisation linearisations

$$h_s^1 = \langle \mathsf{sync}^{1,3}(4, ()):((), 4) \rangle,$$
$$h_s^2 = \langle \mathsf{sync}^{2,3}(5, ()):((), 5) \rangle,$$

corresponding to the two possible synchronisations. Each describes one possibility for all the synchronisations that might happen.

The following definition captures the return events that we would expect to happen given a particular sequence of synchronisations.

Definition 2 Given a history h of the synchronisation object and a maximal synchronisation linearisation h_s, we say that a return event is *anticipated* if it does not appear in h, but the corresponding sync event appears in h_s.

For example, considering the above histories h and h_s^1, the events return.send^1:() and return.receive^3:4 are anticipated: assuming h_s^1 describes the synchronisations that happen, we would expect those return events to occur; if they do not, that is a failure of progress.

The following definition captures our fairness assumption.

Definition 3 Given an execution, we say that an operation execution is *pending* if it has been called but not yet returned. We say that an operation execution is *blocked* in a particular state if it is unable to perform a step, for example because it is trying to acquire a lock held by another thread, or waiting to receive a signal from another thread.

We say that an infinite execution is *fair* if each pending operation execution either (a) eventually returns, (b) is blocked in infinitely many states, or (c) performs infinitely many steps. We consider a system execution that reaches a deadlocked state (where every pending operation is permanently blocked) to be infinite (it is infinite in time), and hence fair (under (b)).

The main reason we require fairness is to rule out executions where a thread obtains a lock, but then is permanently descheduled, which could permanently block other threads. Modern schedulers are fair in this sense.

Definition 4 Let h be a history of the synchronisation object. We say that h is *synchronisation progressible* if for every fair infinite execution (following h) with no new call events, there is a maximal synchronisation linearisation h_s of h such that every anticipated return event ret eventually happens.

For the earlier history h, synchronisation progressibility requires that either return.send^1:() and return.receive^3:4 eventually happen (corresponding to h_s^1), or return.send^2:() and return.receive^3:5 eventually happen (corresponding to h_s^2): one of the synchronisations should happen, and then the relevant operations should return.

On the other hand, for the history $\langle$call.send^1(4)$\rangle$, the only maximal synchronisation linearisation is $\langle\rangle$, for which there are no anticipated returns, and so synchronisation progressibility is trivially satisfied: it is fine for the send to get stuck in this case, since there is no receive with which it could synchronise.

15

3 Comparing synchronisation linearisation and standard linearisation

In this section we describe the relationship between synchronisation linearisation and standard linearisation.

It is clear that synchronisation linearisation is equivalent to standard linearisation in the (rather trivial) case that no operations actually synchronise, so each operation of the synchronisation specification object corresponds to a single operation of the concurrent object. Put another way: standard linearisation is an instance of synchronisation linearisation with just unary synchronisations.

However, linearisability and synchronisation linearisability are not equivalent in general. For example, consider the a synchronous channel from Section 2, where synchronisation linearisation is captured by SyncChanSpec. We show that there is no corresponding linearisability specification Spec, i.e. such that for every concurrent history h, h is synchronisation linearisable with respect to SyncChanSpec if and only if h is linearisable with respect to Spec. Suppose (for a contradiction) otherwise. Consider the history

$$h = \langle \mathsf{call.send}^1(3), \mathsf{call.receive}^2(), \mathsf{return.send}^1:(), \mathsf{return.receive}^2():3 \rangle.$$

This is synchronisation linearisable with respect to SyncChanSpec. By the assumption, it must also be linearisable with respect to Spec; so there must be a legal history h_s of Spec such that h_s is a linearisation of h. Without loss of generality, suppose the send in h_s occurs before the receive, i.e.

$$h_s = \langle \mathsf{send}^1(3):(), \mathsf{receive}^2():3 \rangle.$$

But the history

$$h' = \langle \mathsf{call.send}^1(3), \mathsf{return.send}^1:(), \mathsf{call.receive}^2(), \mathsf{return.receive}^2():3 \rangle$$

is also linearised by h_s, so h' is linearisable with respect to Spec. But then the assumption would imply that h' is synchronisation linearisable with respect to SyncChanSpec. This is clearly false, because the operations do not overlap. Hence no such linearisability specification Spec exists.

3.1 Two-step linearisability

We now show that binary heterogeneous synchronisation linearisability corresponds to a small adaptation of linearisability, where one of the operations

16

on the synchronisation object corresponds to *two* operations of the linearisability specification object. We define below what we mean by this. We prove the correspondence in the next subsection. We generalise to synchronisations of more than two threads, and to the homogeneous case in Section 3.3.

Given a synchronisation object with signature

```
class SyncObj{
  def op₁(x₁: A₁): B₁
  def op₂(x₂: A₂): B₂
}
```

we will consider a linearisability specification object with signature

```
class TwoStepLinSpec{
  def op₁(x₁: A₁): Unit
  def op̄₁(): B₁
  def op₂(x₂: A₂): B₂
}
```

The idea is that the operation op_1 on $\mathsf{SyncObj}$ will be linearised by the composition of the two operations op_1 and $\overline{\mathsf{op}}_1$ of $\mathsf{TwoStepLinSpec}$; but operation op_2 on $\mathsf{SyncObj}$ will be linearised by just the operation op_2 of $\mathsf{TwoStepLinSpec}$, as before. We call such an object a *two-step linearisability specification object*.

We describe how to define a two-step linearisability specification object corresponding to a given synchronisation specification object in the following subsections. First, we define our notion of two-step linearisation with respect to such a two-step linearisability specification object. In the definitions below, we describe just the differences from standard linearisation, to avoid repetition.

We define a history h_s of such a two-step specification object much as in Section 2.2, with the addition that for each event $\overline{\mathsf{op}}_1^i():y$ in h_s, we require that there is an earlier event $\mathsf{op}_1^i(x):()$ in h_s with the same execution identity; other than in this regard, execution identities are not repeated in h_s.

Let h be a complete concurrent history of a synchronisation object, and let h_s be a legal history of a two-step specification object. We say that h_s *corresponds* to h if:

- For every $\mathsf{call}.\mathsf{op}_1^i(x)$ and $\mathsf{return}.\mathsf{op}_1^i:y$ in h, h_s contains $\mathsf{op}_1^i(x):()$ and $\overline{\mathsf{op}}_1^i():y$; and vice versa;

- For every $\mathsf{call}.\mathsf{op}_2^i(x)$ and $\mathsf{return}.\mathsf{op}_2^i:y$ in h, h_s contains $\mathsf{op}_2^i(x):y$; and vice versa.

We say that h_s is a *two-step linearisation* of h if there is some way of interleaving the two histories such that

17

- Each $\mathsf{op}_1^i(x){:}()$ and $\overline{\mathsf{op}}_1^i(){:}y$ occur between $\mathsf{call}.\mathsf{op}_1^i(x)$ and $\mathsf{return}.\mathsf{op}_1^i{:}y$, in that order;

- Each $\mathsf{op}_2^i(x){:}y$ occurs between $\mathsf{call}.\mathsf{op}_2^i(x)$ and $\mathsf{return}.\mathsf{op}_2^i{:}y$.

For example, consider a synchronous channel, with send corresponding to op_1, and $\mathsf{receive}$ corresponding to op_2. Then the following would be an interleaving of histories of the channel and a two-step linearisation specification object.

$$\langle \mathsf{call}.\mathsf{send}^1(3),\ \mathsf{send}^1(3){:}(),\ \mathsf{call}.\mathsf{receive}^2(),\ \mathsf{receive}^2(){:}3,$$
$$\overline{\mathsf{send}}^1(){:}(),\ \mathsf{return}.\mathsf{send}^1{:}(),\ \mathsf{return}.\mathsf{receive}^2{:}3\rangle.$$

This is represented by the following timeline, where the horizontal lines and the labels above represent the execution of operations on the channel, and the "×"s and the labels below represent the corresponding operations of the specification object.

The definition of two-step linearisability of a synchronistion object then follows from this definition of two-step linearisation of complete histories, precisely as in Section 2.2.

3.2 Proving the relationship

We now prove the relationship between synchronisation linearisation and two-step linearisation. Consider a synchonisation specification object $\mathsf{SyncSpec}$. We build a corresponding two-step linearisation specification object $\mathsf{TwoStepLinSpec}$ such that synchronisation linearisation with respect to $\mathsf{SyncSpec}$ is equivalent to two-step linearisation with respect to $\mathsf{TwoStepLinSpec}$. The definition we choose is not the simplest possible, but it is convenient for the testing framework we use in Section 5.

The definition of $\mathsf{TwoStepLinSpec}$ is below. We assume that each thread has an identity in some range $[0 .. \mathsf{NumThreads})$. For simplicity, we arrange for this identity to be included in the call events for operations op_1 and $\overline{\mathsf{op}}_1$.

The object $\mathsf{TwoStepLinSpec}$ encodes the automaton on the right. It requires that corresponding executions of op_1 and op_2 occur consecutively, which means that corresponding executions of op_1 and op_2 on the syn-

chronisation object are linearised successively. However, it allows the corresponding $\overline{op}_1$ to be later (but before the next operation execution by the same thread). It uses an array returns, indexed by thread identities, to record the value that should be returned by an $\overline{op}_1$ execution by each thread. Each execution of op_2 calls SyncSpec.sync to obtain the values that should be returned for synchronisation linearisation; it writes the value for the corresponding $\overline{op}_1$ into returns.

$\overline{op}_1(t)\,[\text{returns}(t) \neq \text{None}]$

Zero

$op_2(x_2)$ $op_1(t, x_1)$ $[\text{returns}(t) = \text{None}]$

One(t, x_1)

```
type ThreadID = Int          // Thread identifiers
val NumThreads: Int = ...     // Number of threads
trait State
case object Zero extends State
case class One(t: ThreadID, x₁: A₁) extends State

object TwoStepLinSpec{
  private var state: State = Zero
  private val returns = new Array[Option[B₁]](NumThreads)
  for(t <− 0 until NumThreads) returns(t) = None
  def op₁(t: ThreadID, x₁: A₁): Unit = {
    require(state == Zero && returns(t) == None); state = One(t, x₁); ()
  }
  def op₂(x₂: A₂): B₂ = {
    require(state.isInstanceOf[One]); val One(t, x₁) = state
    val (y₁, y₂) = SyncSpec.sync(x₁, x₂); returns(t) = Some(y₁); state = Zero; y₂
  }
  def op̄₁(t: ThreadID): B₁ = {
    require(state == Zero && returns(t).isInstanceOf[Some[B₁]])
    val Some(y₁) = returns(t); returns(t) = None; y₁
  }
}
```

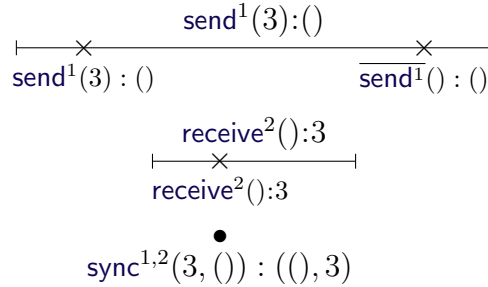The following lemma identifies important properties of TwoStepLinSpec. It follows immediately from the definition.

Lemma 5 Within any legal history of TwoStepLinSpec, events op_1 and op_2 alternate. Let $op_1^{i_1}(t, x_1):()$ and $op_2^{i_2}(x_2):y_2$ be a consecutive pair of such events. Then op_2 makes a call SyncSpec.sync(x_1, x_2) obtaining result (y_1, y_2). The next event for thread t (if any) will be $\overline{op}_1^{i_1}(t):y_1$; and this will be later in the history than $op_2^{i_2}(x_2):y_2$. Further, the corresponding history of events $\text{sync}^{i_1,i_2}(x_1, x_2):(y_1, y_2)$ is a legal history of SyncSpec.

Conversely, each history with events ordered in this way will be a legal history of TwoStepLinSpec if the corresponding history of events $\text{sync}^{i_1,i_2}(x_1, x_2):(y_1, y_2)$ is a legal history of SyncSpec.

The following proposition reduces synchronisation linearisability to two-step linearisability.

Proposition 6 Let SyncObj be a binary heterogeneous synchronisation object, SyncSpec a corresponding synchronisation specification object, and let TwoStepLinSpec be built from SyncSpec as above. Then each history h of SyncObj is two-step linearisable with respect to TwoStepLinSpec if and only if it is synchronisation linearisable with respect to SyncSpec. Hence SyncObj is two-step linearisable with respect to TwoStepLinSpec if and only if it is synchronisation linearisable with respect to SyncSpec.

Proof: ($\Rightarrow$). Let h be a concurrent history of SyncObj that is two-step linearisable with respect to TwoStepLinSpec. By assumption, there is an extension h' of h, and a legal history h_s of TwoStepLinSpec such that h_s is a two-step linearisation of $h'' = complete(h')$. Build a history h'_s of SyncSpec by replacing each consecutive pair $\text{op}_1^{i_1}(x_1):()$, $\text{op}_2^{i_2}(x_2):y_2$ in h_s by the event $\text{sync}^{i_1,i_2}(x_1, x_2):(y_1, y_2)$, where y_1 is the value returned by the corresponding $\overline{\text{op}}_1^{i_1}()$. This is illustrated by the example timeline below, where h'' is represented by the horizontal lines and the labels above; h_s is represented by the "$\times$"s and the labels below; and h'_s is represented by the "$\bullet$" and the label below.

$$\text{send}^1(3):()$$

$$\text{send}^1(3):() \qquad \overline{\text{send}^1}():()$$

$$\text{receive}^2():3$$

$$\text{receive}^2():3$$

$$\bullet$$
$$\text{sync}^{1,2}(3,()):((),3)$$

The history h'_s is legal for SyncSpec by Lemma 5. It is possible to interleave h'' and h'_s by placing each event $\text{sync}^{i_1,i_2}(x_1, x_2):(y_1, y_2)$ in the same place as the corresponding event $\text{op}_2^{i_2}(x_2):y_2$ in the interleaving of h'' and h_s; by construction, this is between $\text{call.op}_1^{i_1}(x_1)$ and $\text{return.op}_1^{i_1}:y_1$, and between $\text{call.op}_2^{i_2}(x_2)$ and $\text{return.op}_2^{i_2}:y_2$. Hence h_s is a synchronisation linearisation of h''; and so h is synchronisation-linearisable.

($\Leftarrow$). Let h be a history of SyncObj that is synchronisation linearisable with respect to SyncSpec. By assumption, there is an extension h' of h, and

a legal history h_s of SyncSpec such that h_s is a synchronisation linearisation of $h'' = complete(h')$. Build a history h'_s of TwoStepLinSpec by replacing each event $\mathsf{sync}^{i_1,i_2}(x_1, x_2){:}(y_1, y_2)$ in h_s by the three consecutive events $\mathsf{op}_1^{i_1}(x_1){:}()$, $\mathsf{op}_2^{i_2}(x_2){:}y_2$, $\overline{\mathsf{op}}_1^{i_1}(){:}y_1$.

The history h'_s is legal for TwoStepLinSpec by Lemma 5. It is possible to interleave h'' and h'_s by placing each triple $\mathsf{op}_1^{i_1}(x_1){:}()$, $\mathsf{op}_2^{i_2}(x_2){:}y_2$, $\overline{\mathsf{op}}_1^{i_1}(){:}y_1$ in the same place as the corresponding event $\mathsf{sync}^{i_1,i_2}(x_1, x_2){:}(y_1, y_2)$ in the interleaving of h'' and h_s; by construction, each $\mathsf{op}_1^{i_1}(x_1){:}()$ and $\overline{\mathsf{op}}_1^{i_1}(){:}y_1$ are between $\mathsf{call.op}_1^{i_1}(x_1)$ and $\mathsf{return.op}_1^{i_1}{:}y_1$; and each $\mathsf{op}_2^{i_2}(x_2){:}y_2$ is between $\mathsf{call.op}_2^{i_2}(x_2)$ and $\mathsf{return.op}_2^{i_2}{:}y_2$. Hence h_s is a two-step linearisation of h''; and so h is two-step-linearisable. $\qquad\square$

The two-step linearisation specification object can often be significantly simplified from the template definition above. Here is such a specification object for a synchronous channel.

```
object SyncChanTwoStepLinSpec{
  private var state = 0         // Takes values 0, 1, cyclically
  private var threadID = −1    // Current sender's thread ID when state = 1
  private var value: A = _      // The current value being sent when state = One
  private val canReturn =       // which senders can return?
    new Array[Boolean](NumThreads)
  def send(t: ThreadID, x: A): Unit = {
    require(state == 0 && !canReturn(t)); value = x; threadID = t; state = 1 }
  def receive(u: Unit): A = {
    require(state == 1); canReturn(threadID) = true; state = 0; value }
  def send(t: ThreadID): Unit = {
    require(state == 0 && canReturn(t)); canReturn(t) = false }
}
```

3.3 Generalisations

The results of the previous subsections carry across to non-binary, fixed arity synchronisations, in a straightforward way. For a k-ary synchronisation between distinct operations $\mathsf{op}_1, \ldots, \mathsf{op}_k$, the corresponding two-step linearisation specification object has $2k-1$ operations, $\mathsf{op}_1, \ldots, \mathsf{op}_k, \overline{\mathsf{op}}_1, \ldots, \overline{\mathsf{op}}_{k-1}$. The definition of two-step linearisation is then the obvious adaptation of the binary case: each operation op_i of the synchronisation object is linearised by the composition of op_i and $\overline{\mathsf{op}}_i$ of the specification object, for $i = 1, \ldots, k-1$.

The construction of the previous subsection is easily adapted to the case of k-way synchronisations for $k > 2$. The specification object encodes an automaton with k states. The figure below gives the automaton in the case $k = 4$.

$$\overline{\mathsf{op}}_i(t)[\mathsf{returns}_i(t) \neq \mathsf{None}]$$
$$(i = 1, 2, 3, 4)$$
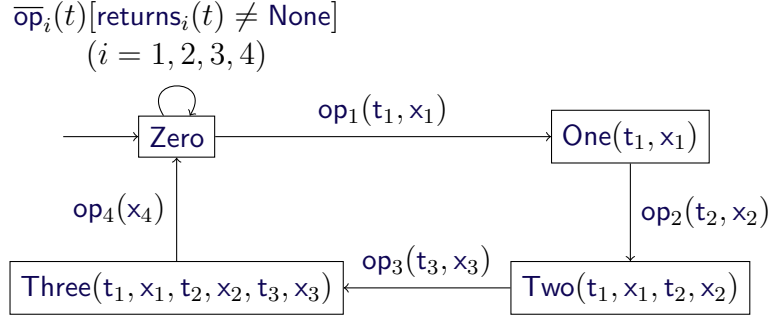
The final op operation, op_4 in the above figure, applies the sync method of the synchronisation specification object to the parameters $\mathsf{x}_1, \ldots, \mathsf{x}_k$ to obtain the results $\mathsf{y}_1, \ldots, \mathsf{y}_k$; it stores the first $k - 1$ in appropriate $\mathsf{returns}_i$ arrays, and returns y_k itself. In the case $k = 4$, it has definition:

```
def op₄(x₄: A₄): B₄ = {
  require(state.isInstanceOf[Three]); val Three(t₁, x₁, t₂, x₂, t₃, x₃) = state
  val (y₁, y₂, y₃, y₄) = SyncSpec.sync(x₁, x₂, x₃, x₄)
  returns₁(t₁) = Some(y₁); returns₂(t₂) = Some(y₂); returns₃(t₃) = Some(y₃)
  state = Zero; y₄
}
```

Each $\overline{\mathsf{op}}_i$ operation retrieves the result from the corresponding $\mathsf{returns}_i$ array.

We now consider the homogeneous case. For simplicity, we describe the binary case; synchronisations of more than two operation executions are handled similarly. Suppose we have a synchronisation object with a single operation **def** op(x: A): B. All executions of op have to be treated similarly, so we associate *each* with two operations op and $\overline{\mathsf{op}}$ of the specification object. The specification object is below, and encodes the automaton on the right. The second execution of op in any synchronisation (from the One state of the automaton) writes the results of the execution into the returns array. Each execution of $\overline{\mathsf{op}}$ returns the stored value.

```scala
class TwoStepLinSpec{
  private var state: State = Zero
  private val returns = new Array[Option[B₁]](NumThreads)
  for(t <− 0 until NumThreads) returns(t) = None
  def op(t: ThreadID, x: A): Unit = {
    require(returns(t) == None)
    state match{
      case Zero => state = One(t, x)
      case One(t₁, x₁) =>
        val (y₁, y₂) = SyncSpec.sync(x₁, x)
        returns(t₁) = Some(y₁); returns(t) = Some(y₂); state = Zero
    }
  }
  def op(t: ThreadID): B = {
    require(state.isInstanceOf[Zero] && returns(t).isInstanceOf[Some])
    val Some(y) = returns(t); returns(t) = None; y
  }
}
```


Recall that some operations have different modes of synchronisations: different executions of the operation may have synchronisations with different arities. For example, in a timeout channel, an execution of the send and receive operations may synchronise with an execution of the other operation, or may timeout corresponding to a unary synchronisation.

Figure 3 gives the automaton for a timeout channel, where we treat send as corresponding to op_1 (we omit concrete code in the interests of brevity). The automaton is similar to that for a standard channel. The receive operation can happen from either state: if it happens from the One state, then a synchronisation has occurred and the execution returns a value of the form Some(x); but if it happens from the Zero state, there has been no corresponding send, and so the execution returns None, indicating a timeout. Likewise, the send operation can happen from either state; if it happens from the Zero state, then a synchronisation has occurred and the execution returns true; but if it happens from the One state, there has been no corresponding receive, and so the execution returns false, indicating a timeout.

We now consider stateful specification objects. In general, we can simply augment the Zero and One states of the automaton to include the state of the specification object. Note that different transitions may be available based upon that state. However, it can be clearer and simpler to introduce different named states into the automaton.

Figure 4 gives the specification automaton for a closeable channel. The automaton has an additional state, Closed, corresponding to the channel being

Figure 3: Automaton for capturing two-step linearisation for a timeout channel.


Figure 4: Automata for capturing two-step linearisation for a closeable channel.

closed. All executions are unary synchronisations in this state. This automaton represents a simplification over the general approach discussed above: we do not need two separate states after closing. Note that a $\overline{\mathsf{send}}(\mathsf{t})$ transition from the Closed state may succeed if the corresponding synchronisation happened before the channel was closed, in which case $\mathsf{returns}(\mathsf{t}) \neq \mathsf{None}$.

We believe that these techniques can be adapted to all the other synchronisation objects we have considered.

4   Linearisability testing

In the following two sections, we describe techniques for testing whether the implementation of a synchronisation object is synchronisation linearisable with respect to a synchronisation specification object. The techniques are

influenced by the techniques for testing (standard) linearisation [Low16], so we begin by sketching those techniques.

The idea of linearisability testing is as follows. We run several *worker threads*, performing operations (typically chosen randomly) upon the concurrent datatype that we are testing, and logging the calls and returns. More precisely, a thread that performs a particular operation $\mathsf{op}^i(x)$: (1) writes $\mathsf{call.op}^i(x)$ into the log; (2) performs $\mathsf{op}(x)$ on the synchonisation object, obtaining result y, say; (3) writes $\mathsf{return.op}^i{:}y$ into the log. Further, the logging associates each operation execution with an execution $\mathsf{op}(x)$ of the corresponding operation on the specification object.

Once all worker threads have finished, we can use an algorithm to test whether the history is linearisable with respect to the specification object. The algorithm searches for an order to linearise the executions, consistent with what is recorded in the log, and such that the order represents a legal history of the corresponding executions on the specification object. See [Low16] for details of several algorithms.

This process can be repeated many times, so as to generate and analyse many histories. Our experience is that the technique works well. It seems effective at finding bugs, where they exist, typically within a few seconds; for example, we used it to find an error in the concurrent priority queue of [ST05], which we believe had not previously been documented. Further, the technique is easy to use: we have taught it to undergraduate students, who have used it effectively.

Note that this testing concentrates upon the safety property of linearisability, rather than liveness properties such as deadlock-freedom. However, if the concurrent object can deadlock, it is likely that the testing will discover this. Related to this point, it is the responsibility of the tester to define the worker threads in a way that all executions will eventually return, so the threads terminate. For example, consider a partial stack where a pop operation blocks while the stack is empty; here, the tester would need to ensure that threads collectively perform at least as many pushes as pops, to ensure that each pop does eventually return.

Note also that there is potentially a delay between a worker thread writing the call event into the log and actually calling the operation; and likewise there is potentially a delay between the operation returning and the thread writing the return event into the log. However, these delays do not generate false errors: if a history without such delays is linearisable, then so is a corresponding history with delays. We believe that it is essential that the technique does not give false errors: an error reported by testing should represent a real error; testing of a correct implementation should be able to run unsupervised, maybe for a long time. Further, our experience is that the

delays do not prevent the detection of bugs when they exist (although might require performing the test more times). This means that a failure to find any bugs, after a large number of tests, can give us good confidence in the correctness of the concurrent datatype.

5   Hacking the linearisablity framework

In this section we investigate how to use the existing linearisation testing framework for testing synchronisation linearisation, using the ideas of Section 3. This is not a use for which the framework was intended, so we consider it a hack. However, it has the advantage of not requiring the implementation of any new algorithms. (We do not consider progressibility in this section.)

Recall, from the introduction of Section 3, that a straightforward approach won't work. Instead we adapt the idea of two-step linearisation from later in that section. We start by considering the case of binary heterogeneous synchronisation. We aim to obtain a log history that can be tested for (standard) linearisability against TwoStepLinSpec.

As with standard linearisability testing, we run several worker threads, calling operations on the synchronisation object, and logging the calls and returns.

- A thread t_1 that performs the concrete operation $\mathsf{op}_1(x_1)$: (1) writes $\mathsf{call}.\mathsf{op}_1^{i_1}(x_1)$ into the log, associating it with a corresponding execution $\mathsf{op}_1(t_1, x_1)$ on the specification object; (2) performs $\mathsf{op}_1(x_1)$ on the synchonisation object, obtaining result y_1, say; (3) writes $\mathsf{return}.\mathsf{op}_1^{i_1}{:}()$ into the log; (4) writes $\mathsf{call}.\overline{\mathsf{op}}_1^{i_1}()$ into the log, associating it with a corresponding execution $\overline{\mathsf{op}}_1(t)$ on the specification object; (5) writes $\mathsf{return}.\overline{\mathsf{op}}_1^{i_1}{:}y_1$ into the log.

- A thread t_2 that performs operation op_2, acts as for standard linearisability testing. It: (1) writes $\mathsf{call}.\mathsf{op}_2^{i_2}(x_2)$ into the log, associating it with a corresponding execution $\mathsf{op}_2(x_2)$ on the specification object; (2) performs $\mathsf{op}_2(x_2)$ on the synchonisation object, obtaining result y_2, say; (3) writes $\mathsf{return}.\mathsf{op}_2^{i_2}{:}y_2$ into the log

The top half of Figure 5 illustrates a possible run, containing a single synchronisation, together with the log history.

The following definition captures that a log history h_l might arise from a concrete history h, using the logging strategy described above.

Definition 7   Let h be a complete history of a binary heterogeneous synchronisation object, and let h_l be a log history for the same object. We say

Figure 5: Illustration of two-step linearisation testing. The operation executions are represented by the horizontal lines with labels above (denoted "h" in Proposition 8). The log entries are represented by the bullets with labels below (denoted "h_l" in Proposition 8). Linearisation points are represented by crosses with labels below: the penultimate row, labelled "h_s", is a synchronisation linearisation; the bottom row, labelled "h_{2s}", is a linearisation of the two-step synchronisation object. Execution identifiers and null arguments and returns are omitted, for clarity.

that the two histories *correspond* if there is some way of interleaving them such that

- Each $\mathsf{call}.\mathsf{op}_1^{i_1}(x_1)$, from h_l, precedes the call and return of $\mathsf{op}_1^{i_1}(x_1){:}y_1$ from h, which precedes $\mathsf{return}.\mathsf{op}_1^{i_1}{:}()$, $\mathsf{call}.\overline{\mathsf{op}}_1^{i_1}()$ and $\mathsf{return}.\overline{\mathsf{op}}_1^{i_1}{:}y_1$, from h_l, in that order.

- Each $\mathsf{call}.\mathsf{op}_2^{i_2}(x_2)$, from h_l, precedes the call and return of $\mathsf{op}_2^{i_2}(x_2){:}y_2$ from h, which precedes $\mathsf{return}.\mathsf{op}_2^{i_2}{:}y_2$, from h_l.

As with standard linearisation, the tester needs to define the worker threads so that all executions will eventually return, i.e. that each will be able to synchronise. For a binary heterogeneous synchronisation with no precondition, we can achieve this by half the threads calling one operation, and the other half calling the other operation (with the same number of calls by each). For a binary homogeneous synchronisation, this approach might not work if every worker does more than one operation: one worker might end up with two operations to perform, when all others have terminated; instead, we arrange for an even number of workers to each perform one operation.

Once all threads have finished, we test whether the log history is linearisable (i.e. standard linearisation) with respect to TwoStepLinSpec from Section 3. Figure 5 gives an example linearisation.

27

Note that we have three related concepts here: (1) synchronisation linearisation of the concrete history of operation executions with respect to SyncSpec; (2) two-step linearisation of the concrete history with respect to TwoStepLinSpec; and (3) linearisation of the log history with respect to TwoStepLinSpec. Proposition 6 shows that the first two of these are equivalent. We need to show that these imply (3), so the technique does not give false errors. (The converse might not hold, because of delays in writing to the log.)

Proposition 8 Let h be a complete history of a binary heterogeneous synchronisation object, and let h_l be a corresponding log history for the same object. Let SyncSpec be a synchronisation specification object, and TwoStepSyncSpec the corresponding two-step synchronisation specification object, constructed as in Section 3.2. Suppose h is synchronisation linearisable with respect to SyncSpec. Then h_l is linearisable with respect to TwoStepSyncSpec.

Proof: Since h is synchronisation linearisable, there is a legal history h_s of SyncSpec such that h_s is a synchronisation linearisation of h. Consider the interleaving of h_s and h, that demonstrates this, and interleave h_l with it, consistent with the interleaving of h and h_l that demonstrates that they correspond. Figure 5 illustrates such an interleaving.

We build a history h_{2s} of TwoStepSyncSpec, and interleave it with h_l as follows. In the interleaving of the previous paragraph, replace each event $\mathsf{sync}^{i_1,i_2}(x_1, x_2){:}(y_1, y_2)$ (from h_s) by consecutive events $\mathsf{op}_1^{i_1}(x_1){:}()$ and $\mathsf{op}_2^{i_2}(x_2){:}y_2$, and add $\overline{\mathsf{op}}_1^{i_1}(){:}y_1$ between $\mathsf{call}.\overline{\mathsf{op}}_1^{i_1}()$ and $\mathsf{return}.\overline{\mathsf{op}}_1^{i_1}{:}y_1$ (from h_l). Again, Figure 5 illustrates such an interleaving. This is a legal history of TwoStepSyncSpec, by Lemma 5. Further, each event of h_{2s} is between the corresponding call and return events of h_l, by construction. Hence h_{2s} is a linearisation of h_l. $\qquad\square$

This approach generalises to non-binary synchronisations, homogeneous synchronisations, and stateful specification objects as in Section 3.3.

Variable-arity synchronisations. It turns out that it is not, in general, possible to capture variable-arity synchronisations using this technique, in particular where the arity of a synchronisation depends upon the relative timing of executions, as opposed to the state of the specification object. This is a result of two things: that the logging of operations, in particular the $\overline{\mathsf{op}}_1$, can be arbitrarily delayed; and that it can be nondeterministic whether or not two executions synchronise, which is at odds with the fact that each operation on the specification object needs to be deterministic.
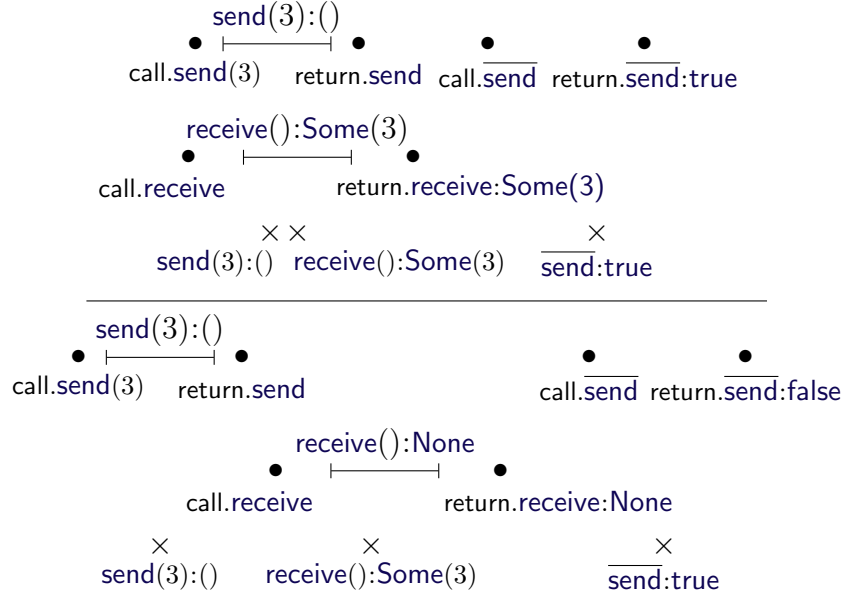
Figure 6: Figure showing why two-step linearisation cannot be used for a timeout channel. Conventions are as in Figure 5.

To illustrate this point, consider a timeout channel. Without loss of generality, let the send operation correspond to op_1, and the receive operation correspond to op_2.

The top-half of Figure 6 gives a timeline illustrating a successful send(3) and receive. Note that this corresponds to the history

$$\langle \text{send}(3){:}(), \ \text{receive}(){:}\text{Some}(3), \ \overline{\text{send}}(3){:}\text{true} \rangle$$

of the specification object.

The bottom-half side of Figure 6 gives a timeline illustrating an unsuccessful send(3) and receive, and where the logging of $\overline{\text{send}}$ is delayed. None of the executions overlap, so they must necessarily be linearised in the same order as in the previous history. The specification object is deterministic, so the operations must return the same results as in the previous history. But, in the cases of receive and $\overline{\text{send}}$, those returned values do not agree with the corresponding values in the log history. Hence the history would be flagged as an error, despite being valid.

The difference between this situation and the discussion in Section 3.3 concerns the fact that the logging of operations, in particular the $\overline{\text{send}}$, can be arbitrarily delayed. However, in the earlier section we allowed the $\overline{\text{send}}$

anywhere within the corresponding concrete operation. This means that a history like in the bottom-half of Figure 6 could be linearised by the history

$$\langle \mathsf{send}(3){:}(),\ \overline{\mathsf{send}}(3){:}\mathsf{false},\ \mathsf{receive}(){:}\mathsf{None} \rangle$$

of the two-step specification object, where the operations take place in a different order than for Figure 6; this is consistent with a deterministic specification object.

A similar problem arises with a timeout exchanger.

We have investigated an alternative approach, which involves worker threads adapting their logging behaviour based on the outcome of their operations. For the timeout channel:

- A thread that sends a value x: (1) writes $\mathsf{call.send}^{i_1}(x)$ into the log; (2) performs $\mathsf{send}(x)$ on the channel; (3) writes $\mathsf{return.send}^{i_1}{:}()$ into the log; (4) if the send is successful, associates the log entries with an operation $\mathsf{send}(x)$ on the specification object, and otherwise associates them with an operation $\mathsf{sendFail}(x)$; (5) if the send is successful, writes $\mathsf{call.\overline{send}}^{i_1}()$ and $\mathsf{return.\overline{send}}_1^{i_1}{:}()$ into the log, associating them with an operation $\overline{\mathsf{send}}()$ on the specification object (and otherwise does nothing).

- A thread that performs a receive: (1) writes $\mathsf{call.receive}^{i_2}$ into the log; (2) performs $\mathsf{receive}$ on the channel, receiving result r, say; (3) writes $\mathsf{return.receive}^{i_1}{:}r$ into the log; (4) if the receive was successful, associates the log entries with an operation $\mathsf{receive}$ on the specification, and otherwise associates them with an operation $\mathsf{receiveFail}$.

The specification object encodes the automaton below.

Thus a successful synchronisation is linearised by a sequence $\mathsf{send}(x)$, $\mathsf{receive:Some}(x)$, $\overline{\mathsf{send}}$, with the former two events consecutive, as for other binary heterogeneous synchronisations. Unsuccessful sends and receives are each linearised by a single event.

A similar technique can be used for a timeout exchanger. We consider this approach convoluted, and we do not advocate it. We include it only for completeness.

6 Direct testing of synchronisation linearisation and progressability

We now consider how to test for synchronisation linearisation more directly. We perform logging precisely as for standard linearisation: a thread that performs a particular operation $op^i(x)$: (1) writes $call.op^i(x)$ into the log; (2) performs $op(x)$ on the synchonisation object, obtaining result y, say; (3) writes $return.op^i{:}y$ into the log.

When testing for synchronisation linearisation, we again make it the responsibility of the tester to define the worker threads in a way that ensures that all operation executions will be able to synchronise, so all threads will eventually terminate.

When testing for progress, we remove the requirement on the tester to ensure that all operation executions can synchronise. Indeed, in some cases, in order to find failures of progress, it is necessary that not all executions can synchronise: we have examples of incorrect synchronisation objects where (for example) if there are two executions of op_1 and one of op_2, then it's possible that *neither* execution of op_1 returns, signifying a failure of progressability; but if there were a second execution of op_2, it would unblock both executions of op_1, so all executions would return, and the failure of progressability would be missed.

Instead, we run threads performing operations, typically chosen at random; and after a suitable duration, we interrupt any threads that have not yet returned. The duration before the interrupts needs to be chosen so that it is highly likely that any threads that have not returned really are stuck: otherwise this approach it likely to produce false positives. Our informal experiments suggest that a duration of 100ms is appropriate (at least, on the architecture we were using): we have not observed any false positives with this duration, but did with a shorter duration of 80ms. This delay does significantly increase the time that a given number of runs will take.

In the remainder of this section we consider algorithms for determining if the resulting log history is synchronisation linearisable, and whether it is synchronisation progressable. In Section 6.1 we present a general algorithm for this problem, based on depth-first search. We then consider the complexity of this problem. We show, in Section 6.2, that, in the case of a stateful synchronisation object, the problem of deciding whether a history is synchronisation linearisable is NP-complete in general. However, we show that in the case of binary synchronisations with a stateless specification object the problem can be solved in polynomial time: we consider the heterogeneous case in Section 6.3, and the homogeneous case in Section 6.4. Nevertheless,

31

in Section 6.5 we show that for synchronisations of three or more executions, the problem is again NP-complete, even in the stateless case.

6.1 The general case

We describe an algorithm for deciding whether a given complete history h is synchronisation linearisable with respect to a given synchronisation specification object. We transform the problem into a graph-search algorithm as follows.

We define a search graph, where each node is a *configuration* comprising:

- An index i into the log;

- A set *pending* of operation executions that were called in the first i events of the log and that have not yet been linearised;

- A set *linearised* of operation executions that were called in the first i events of the log and that have been linearised, but have not yet returned;

- The state *spec* of the specification object after the synchronisations linearised so far.

From such a configuration, there are edges to configurations as follows:

Synchronisation. If some set of executions in *pending* can synchronise, giving results compatible with *spec*, then there is an edge to a configuration where the synchronising executions are moved into *linearised*, and *spec* is updated corresponding to the synchronisation;

Call. If the next event in the log is a call event, then there is an edge where that event is added to *pending*, and i is advanced;

Return. If the next event in the log is a return event, and the corresponding execution is in *linearised*, then there is an edge where that execution is removed from *linearised*, and i is advanced.

The initial configuration has i at the start of the log, *pending* and *linearised* empty, and *spec* the initial state of the specification object. Target configurations have i at the end of the log, and *pending* and *linearised* empty.

Any path from the initial configuration to a target configuration clearly represents an interleaving of a history of the specification object with h, as required for synchronisation linearisation. We can therefore search this graph using a standard algorithm. Our implementation uses depth-first search.

It is straightforward to adapt the search algorithm to also test for progress. We change the definition of a target configuration to have i at the end of the log, *linearised* empty, and such that no set of executions in *pending* can synchronise: this ensures that we are dealing with a *maximal* synchronisation linearisation.

6.1.1 Partial-order reduction

We have investigated a form of partial-order reduction, which we call *ASAP linearisation*. The idea is that we try to linearise executions *as soon as possible*.

Definition 9 Let h be a complete history of a synchronisation object, and let h_s be a legal history of the corresponding specification object; and consider an interleaving, as required for synchronisation linearisation. We say that the interleaving is an *ASAP interleaving* if every event in h_s appears either: (1) directly after the call event of one of the corresponding executions from h; or (2) directly after another event from h_s.

The following lemma shows that it suffices to consider ASAP interleavings.

Lemma 10 Let h be a complete history of a synchronisation object, and let h_s be a legal history of the corresponding specification object. If h_s is a synchronisation linearisation of h, then there is an ASAP interleaving of them.

Proof: Consider an interleaving of h and h_s, as required for synchronisation linearisation. We transform it into an ASAP interleaving as follows. Working forwards through the interleaving, we move every event of h_s earlier in the interleaving, as far as possible, without it moving past any of the corresponding call events, nor moving past any other event from h_s. This means that subsequently each such event follows either a corresponding call event or another event from h_s.

Note that each event from h_s is still between the call and return events of the corresponding executions. Further, we do not reorder events from h_s so the resulting interleaving is still an interleaving of h and h_s.

Thus the resulting interleaving is an ASAP interleaving. □

Our approach, then, is to trim the search graph by removing synchronisation edges that do not correspond to an ASAP linearisation: after a call edge, we attempt to linearise a synchronisation corresponding to that call, and then, if successful, to linearise an arbitrary sequence of other synchronisations; but we do not otherwise allow synchronisations.

Our experience is that this tactic is moderately successful. In some cases, it can reduce the total time to check histories by over 30%; although in some cases the gains are smaller, sometimes negligible. The gains seem highest in examples where there can be a reasonably large number of pending executions.

6.2 Complexity

Consider the problem of testing whether a given concurrent history is synchronisation linearisable with respect to a given synchronisation specification object. We show that this problem is NP-complete in general.

We make use of a result from [GK97] concerning the complexity of the corresponding problem for linearisability. Let Variable be a linearisability specification object corresponding to an integer variable with get and set operations. Then the problem of deciding whether a given concurrent history is linearisable with respect to Variable is NP-complete.

Since standard linearisation is a special case of synchronisation linearisation (in the trivial case of unary synchronisations), this immediately implies that deciding synchronisation linearisation is NP-complete. However, even if we restrict to the non-trivial case of binary synchronisations, the result still holds. We consider concurrent synchronisation histories on an object with the following signature, which mimics the behaviour of a variable but via synchronisations.

object VariableSync{
  def op_1(op: String, x: Int): Int
  def op_2(u: Unit): Unit
}

The intention is that op_1("get", x) acts like get(x), and op_1("set", x) acts like set(x) (but returns -1). The op_2 operation does nothing except synchronise with op_1. This can be captured formally by the following synchronisation specification object.

object VariableSyncSpec{
  private var state $= 0$ // *the value of the variable*
  def sync((op, x): (String, Int), u: Unit): (Int, Unit) $=$
    if(op $==$ "get") (state, ()) **else**{ state $=$ x; $(-1, ())$ }
}

Let ConcVariable be a concurrent object that represents an integer variable. Given a history h of ConcVariable, we build a history h' of VariableSync as follows. We replace every call or return of get(x) by (respectively) a call or return of op_1("get", x); and we do similarly with sets. If there are k calls of get

or set in total, we prepend k calls of op$_2$, and append k corresponding returns (in any order). Then it is clear that h is linearisable with respect to Variable if and only if h' is linearisable with respect to VariableSyncSpec. Deciding the former is NP-complete; hence the latter is also.

6.3   The binary heterogeneous stateless case

The result of the previous subsection used a *stateful* specification object. We now consider the *stateless* case. We show that for binary heterogeneous synchronisations, the problem of deciding whether a history is synchronisation linearisable can be decided in quadratic time. We consider the homogeneous case in the next subsection.

So consider a binary heterogeneous synchronisation object, whose specification object is stateless. Note that in this case we do not need to worry about the order of synchronisations: if each individual synchronisation is correct, then any permutation will also be correct from the point of view of the specification object; and we can order the synchronisations in a way that is compatible with the concurrent history. Informally, the idea is to find matching operation executions in the concurrent history that could correspond to a particular synchronisation; we therefore reduce the problem to that of finding a matching in a graph.

Define two complete operation executions to be *compatible* if they could be synchronised, i.e. they overlap and the return values agree with those for the specification object. For n executions of operations this can be calculated in $O(n^2)$.

Consider the bipartite graph where the two sets of nodes are executions of op$_1$ and op$_2$, respectively, and there is an edge between two executions if they are compatible. A synchronisation linearisation then corresponds to a total matching of this graph: given a total matching, we build a synchronisation linearisation by including events sync$^{i_1,i_2}(x_1, x_2){:}(y_1, y_2)$ (in an appropriate order) whenever there is an edge between op$_1^{i_1}(x_1){:}y_1$ and op$_2^{i_2}(x_2){:}y_2$ in the matching; and conversely, each synchronisation linearisation corresponds to a total matching.

Thus we have reduced the problem to that of deciding whether a total matching exits, for which standard algorithms exist. We use the Ford-Fulkerson method [FF56], which runs in time $O(n^2)$.

It is straightforward to extend this to a mix of binary and unary synchronisations, again with a stateless specification object: the executions of unary operations can be considered in isolation.

This approach can be easily extended to also test for progress. It is enough to additionally check that no two pending executions could synchronise.

6.4 The binary homogeneous stateless case

We now consider the case of binary *homogeneous* synchronisations with a stateless specification object. This case is almost identical to the case with heterogeneous synchronisations, except the graph produced is not necessarily bipartite. Thus we have reduced the problem to that of finding a maximum matching in a general graph. This problem can be solved using, for example, the blossom algorithm [Edm65], which runs in time $O(n^4)$.

In fact, our experiments use a simpler algorithm. We attempt to find a matching via a depth-first search: we pick a node n that has not yet been matched, try matching it with some unmatched compatible node n', and recurse on the remainder of the graph; if that recursive search is unsuccessful, we backtrack and try matching n with a different node. We guide this search by the standard heuristic of, at each point, expanding the node n that has fewest unmatched compatible nodes n'.

In our only example of this category, the Exchanger from the Introduction, we can choose the values to be exchanged randomly from a reasonably large range (say size 100). Then we can nearly always find a node n for which there is a unique unmatched compatible node: this means that the algorithm nearly always runs in linear time. We expect that similar techniques could be used in other examples in this category.

6.5 The non-binary stateless case

It turns out that for synchronisations of arity greater than 2, the problem of deciding whether a history is synchronisation linearisable is NP-complete in general, even in the stateless case. We prove this fact by reduction from the following problem, which is known to be NP-complete [Kar72].

Definition 11 The problem of finding a complete matching in a 3-partite hypergraph is as follows: given disjoint finite sets X, Y and Z of the same cardinality, and a set $T \subseteq X \times Y \times Z$, find $U \subseteq T$ such that each member of X, Y and Z is included in precisely one element of T.

Suppose we are given an instance (X, Y, Z, T) of the above problem. We construct a synchronisation specification and a corresponding history h such that h is synchronisation linearisable if and only if a complete matching exists. The synchronisations are between operations as follows:

```
def op₁(x: X): Unit
def op₂(y: Y): Unit
def op₃(z: Z): Unit
```

The synchronisations are specified by:

```
def sync(x: X, y: Y, z: Z): (Unit, Unit, Unit) = {
    require((x, y, z) ∈ T);  ((), (), ())
}
```

The history h starts with calls of $\mathsf{op}_1(x)$ for each $x \in X$, $\mathsf{op}_2(y)$ for each $y \in Y$, and $\mathsf{op}_3(z)$ for each $z \in Z$ (in any order); and then continues with returns of the same executions (in any order). It is clear that any synchronisation linearisation corresponds to a complete matching, i.e. the executions that synchronise correspond to the complete matching U. Hence finding a synchronisation linearisation is NP-complete.

Our implementation for these cases uses a depth-first search to find a matching, very much like in the binary homogeneous case.

6.6   Implementation

We have implemented a testing framework (in Scala), based on the above algorithms, and used it to implement testers for particular synchronisation objects[4]. We consider the framework to be straightforward to use: most of the boilerplate code is encapsulated within the framework; defining a tester for a new synchronisation object takes just a few minutes.

Figure 7 gives a stripped-down tester for a synchronous channel. (The full version can be used to test a number of implementations with the same interface, and replaces the numeric constants by parameters that can be set on the command line.)

The worker function defines a worker thread that performs operations on the channel c. The function also takes parameters representing the thread's identity and a log object. Here, each worker with an even identity performs 20 receive executions: the call log(me, c.receive(), Receive) logs the call, performs the receive, and then logs the return. Similarly, each worker with an odd identity performs 20 send executions of random values. This definition is designed so that an even number of workers with contiguous identities will not deadlock.

SyncChanSpec is the synchronisation specification object from earlier. The way executions synchronise is captured by the function matching. This is a partial function whose domain defines which operation executions can synchronise together, and, in that case, the value each should return: here send(x) and receive can synchronise, giving a result as defined but the synchronisation specification object. (Alternatively, the call to SyncChanSpec.sync can be in-lined.)

---

[4]The implementation is available from ???.

```
object ChanTester extends Tester{
  trait Op   // Representation of operations within the log
  case class Send(x: Int) extends Op
  case object Receive extends Op

  def worker(c: SyncChan[Int])(me: Int, log: HistoryLog[Op]) =
    for(i <− 0 until 20)
      if(me%2 == 0) log(me, c.receive(), Receive)
      else{ val x = Random.nextInt(100); log(me, c.send(x), Send(x)) }

  object SyncChanSpec{
    def sync(x: Int, u: Unit) = ((), x)
  }

  def matching: PartialFunction[(Op,Op), (Any,Any)] = {
    case (Send(x), Receive) => SyncChanSpec.sync(x, ()) // = ((), x)
  }

  /** Do a single test.  Return true if it passes. */
  def doTest(): Boolean = {
    val c = new SyncChan[Int]
    new BinaryStatelessTester[Op](worker(c), 8, matching)()
  }

  def main(args: Array[String]) = {
    var i = 0; while(i < 5000 && doTest()) i += 1
  }
}
```

Figure 7: A simple tester for a synchronous channel.

The function doTest performs a single test. This uses a BinaryStatelessTester
object from the testing framework, which encapsulates the search from Sec-
tion 6.3. Here, the tester runs 8 worker threads, and tests the resulting history
against matching. If a non-synchronisation-linearisable history is recorded, it
displays this for the user. The main function runs doTest either 5000 times or
until an error is found. The tester can be adapted to test for synchronisation
progressibility by passing a timeout duration to the BinaryStatelessTester.

Other classes of testers are similar. In the case of a stateful specification,
the matching function takes the specification object as a parameter, and also
returns the new value of the specification object.

38

Category	Arity	Stateful?	Heterogeneous?
Synchronous channel	2	N	Y
Filter channel	2	N	Y
Men and women	2	N	Y
Exchanger	2	N	N
Counter channel	2	Y	Y
Two families	2	Y	Y
One family	2	Y	N
ABC	3	N	Y
Barrier	n	N	N
Enrollable barrier	$1..n$, 1	Y	N
Timeout channel	2, 1	N	Y
Timeout exchanger	2, 1	N	N
Closeable channel	2, 1	Y	Y
Terminating queue	1, n	Y	N

Figure 8: Example interfaces of synchronisation objects.

7 Experiments

In this section we describe experiments based on our testing framework.

We consider synchronisation objects implementing a number of interfaces, summarised in Figure 8. Most of the interfaces were described in earlier sections (namely synchronous channel, filter channel, exchanger, counter channel, barrier, enrollable barrier, timeout channel, closeable channel, and terminating queue). The *men and women* problem involves two families of threads, known as men and women: each thread wants to pair off with a thread of the other type; each passes in its own identity, and expects to receive back the identity of the thread with which it has paired. In the *two families* problem, there are two families of threads, with n threads of each family; each thread calls an operation n times, and each execution should synchronise with a thread of the opposite family, a different thread each time. In the *one family* problem, there are n threads, each of which calls an operation $n-1$ times, and each time should synchronise with a different thread. The *ABC* problem can be thought of as a ternary version of the men and women problem: there are three types of threads, A, B and C; each synchronisation involves one thread of each type. Finally, the *timeout exchanger* is a timed version of the exchanger: if a thread fails to exchange data with another thread, it can timeout and return an appropriate result.

For each interface, we have implemented a tester using the direct algo-

rithm from Section 6; for most, we have also implemented a tester using the two-step approach from Section 4.

For each interface, we have produced a correct implementation. For most interfaces, we have also implemented one or more faulty versions that fail to achieve either synchronisation linearisation or progressibility. The faulty versions mostly have realistic mistakes: a few are genuine bugs; others are similar bugs we have seen from students.

7.1 Experiments

We describe various experiments below. The purpose of testing is to find bugs. We therefore concentrate on the time taken to find bugs. Questions we want to answer include:

- Which works better, the direct algorithm or the two-step algorithm?

- How should we choose parameters (number of threads to run, number of iterations performed by each thread, etc.) for testing?

- Is this approach effective at finding bugs?

The experiments were performed on a dedicated eight-core machine (two 2.40GHz Intel(R) Xeon(R) E5620 CPUs, with 12GB of RAM, but limited to 4GB of heap space).

In each experiment below, we consider a synchronisation object with a bug that causes a failure of synchronisation linearisation, but does not lead to a deadlock. We performed a number of *runs* of a tester on the synchronisation object. In each run, a particular number of threads performed a particular number of operation calls on the synchronisation object; the relevant algorithm was then used to decide whether the log history was synchronisation linearisable or two-step linearisable.

Each *observation* performed multiple runs until an error was detected, and recorded the time taken. Each observation was performed as a separate operating system process, with the aim of making observations independent, avoiding dependencies caused by, for example, garbage collection, caching behaviour, and just-in-time compilation. Thus each observation was as close as possible to a normal use case.

For each data point in the experiments, we performed 100 observations. We give the average running time and a 95%-confidence interval for that average (following [GBE07]). The number of observations is chosen so as to obtain a reasonably small confidential interval, but avoiding excessively long experiments.

We start with the second of the questions above, how to choose parameters for testing. Each of the graphs in Figures 9 and 10 concerns a particular tester. Each data point represents a particular number p of threads (given in the key), and a particular number of operation executions by each thread (given on the X-axis). The Y-axis gives the time in milliseconds.

The experiments suggest that both types of tester work best with a fairly small number of threads, each performing a fairly small number of operations per run. Some bugs are exhibited only when the number of threads exceeds the arity of a synchronisation: for the ABC problem, two different synchronisations interfere; for the closeable channel, the closing of the channel interferes with a synchronisation. We therefore recommend including enough threads to find such bugs.

These results suggest that the direct algorithms scale better than the two-step approach as the number of threads increases. Figure 11 investigates this further. Each graph considers one type of synchronisation object, with the two plots representing the two testers. Each data point considers a particular number of threads (given on the X-axis). The Y-axis again gives the time in milliseconds. We omit results for the two-step tester where testing sometimes ran out of memory. The results confirm that the two-step approach does not scale well with the number of threads. In each case, the running time explodes at a particular point, and tests often fail. The running time also becomes erratic (as indicated by the wide confidence intervals): some runs take an extremely long time. By contrast, the direct algorithms scale well.

We believe the reason the two-step approach scales poorly is as follows. The linearisation tester tries to find a linearisation order for operation executions, via a depth-first search. At each step, it picks a particular operation execution to try to linearise next. If it picks wrongly, it might have to consider many nodes of the search graph before backtracking. This can be the case with two-step testing, because if it picks the wrong op_1 to try to linearise, it will only discover this fact when it reaches the corresponding $\overline{\mathsf{op}}_1$ and finds the wrong value is returned. This latter event might be much later in the history. To reach it, the tester has to consider many possibilities for ordering other operation executions.

Figure 12 gives times to find various bugs with the two testers. Based on the earlier experiments, in most cases we ran four threads, each executing four operations; for the ABC testers, we ran six threads, each executing four operations; for the (untimed) exchanger, we ran eight threads, each performing one operation; for the two-families object, we ran four threads (two from each family), each performing two operations; and for the one-family object, we ran four threads, each performing three operations. The table gives average times in milliseconds to detect the bug, with 95%-confidence
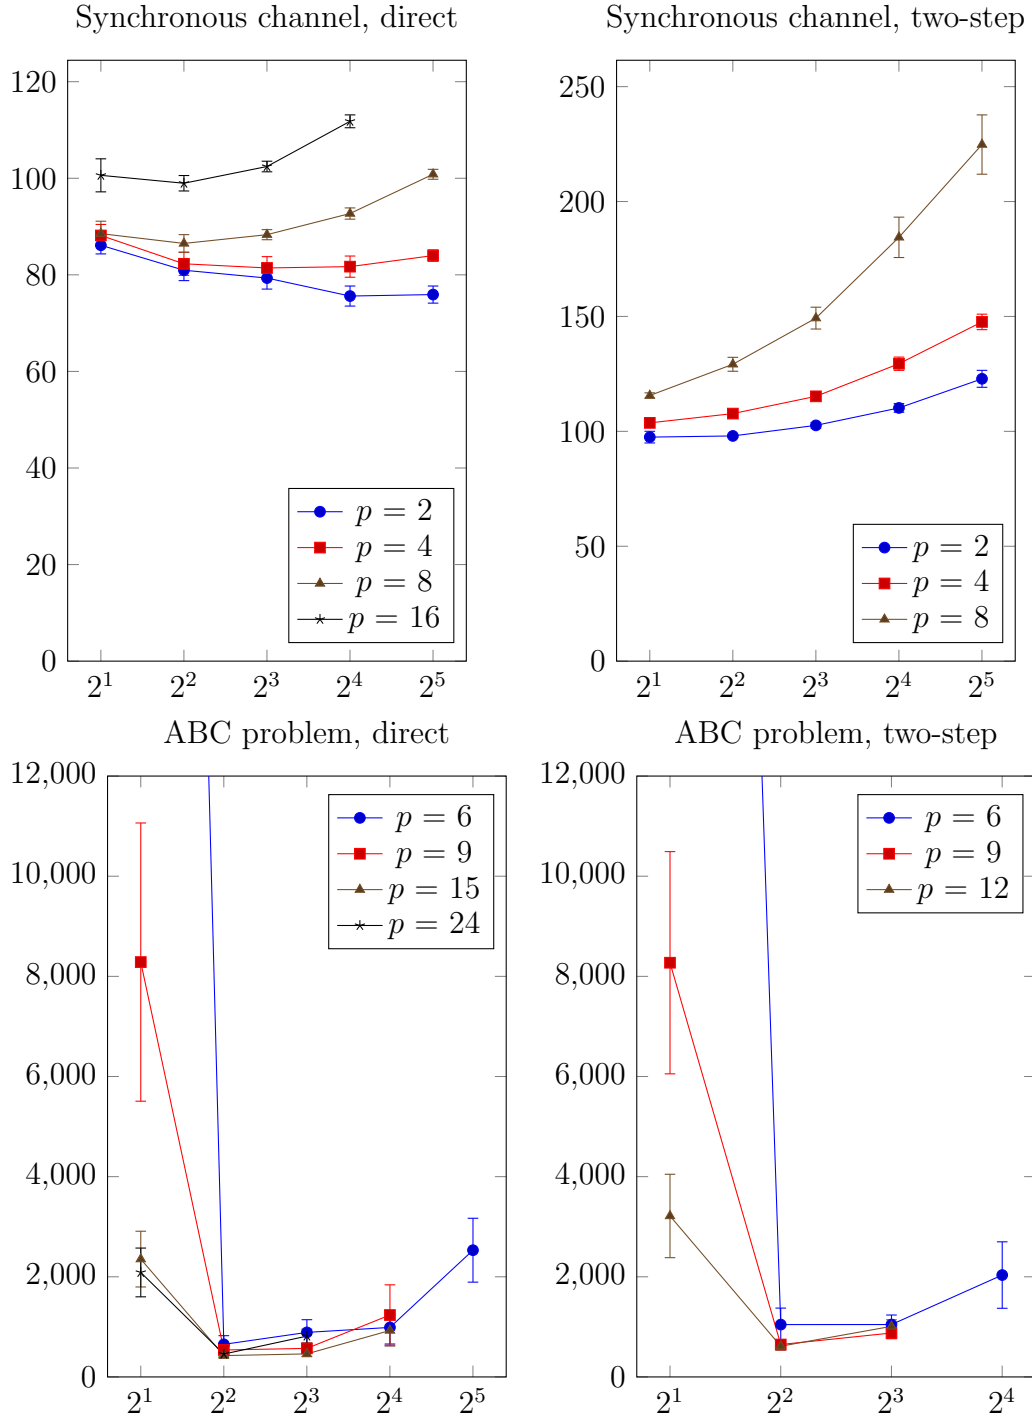
Figure 9: Effect of choices of parameters for testers for a synchronous channel and the ABC problem.
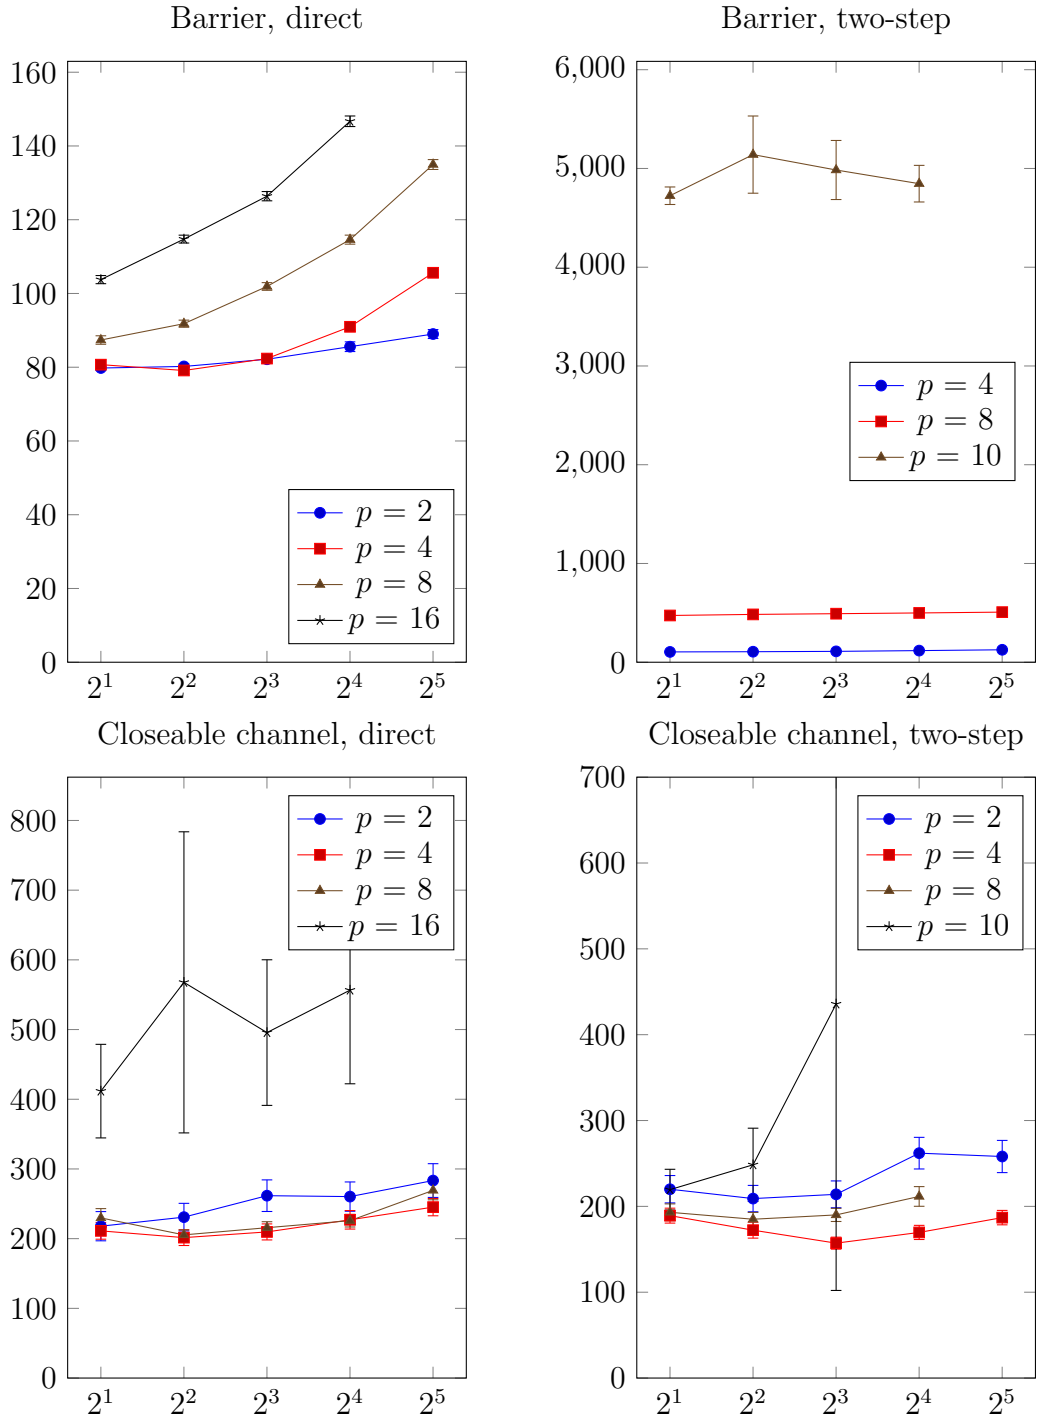
Figure 10: Effect of choices of parameters for testers for a barrier and a closeable channel.

Figure 11: The effect of increasing the number of threads on the time to find bugs.

Synchronisation object	Direct	Two-step
Synchronous channel	85 ± 3	109 ± 2
Filter channel	—	—
Men and women	75 ± 1	108 ± 4
Exchanger	73 ± 1	511 ± 25
Counter channel	94 ± 2	106 ± 1
Two families	299 ± 34	267 ± 25
One family	336 ± 27	—
ABC	717 ± 225	928 ± 216
Barrier	80 ± 1	106 ± 1
Enrollable barrier	132 ± 5	149 ± 6
Timeout channel	121 ± 5	135 ± 5
Timeout exchanger	288 ± 64	231 ± 18
Closeable channel	191 ± 11	173 ± 9
Terminating queue	103 ± 1	105 ± 1

??? Add one family, filter channel?

Figure 12: Times to find bugs.

intervals.

All the testers work well, with the average time to find each bug below one second. Of course, other bugs might be harder to find, because they are triggered on fewer runs. However, our results do suggest that our techniques are effective at finding most bugs.

In most cases, the direct tester is faster than the two-step one; and the two-step tester is never significantly faster. We therefore recommend the direct algorithms. This approach has an additional advantage: our experience is that it is easier to create the testing program based on these algorithms, whereas using the two-step approach involves designing and encoding the appropriate automaton, which can be somewhat tricky.

We now consider synchronisation progressibility. Our experience is that if a synchronisation object does not satisfy progressibility, then this can lead to a total deadlock. Thus, in most cases, testing for synchronisation linearisability will also detect progressibility bugs. However, this is no guarantee.

Add experiments?

8   Conclusions

Short errors easier to interpret.

Model checking.

Spurious wakeups.
Ease of writing testers.

References

[Edm65]   Jack Edmonds. Paths, trees, and flowers. *Canadian Journal of Mathematics*, 17:449–467, 1965.

[FF56]   Lestor R. Ford, Jr. and D. R. Fulkerson. Maximal flow through a network. *Canadian Journal of Mathematics*, 8:399–404, 1956.

[GBE07]   Andy Georges, Dries Buytaert, and Lieven Eeckhout. Statistically rigorous Java performance evaluation. In *Proceedings of OOPSLA '07*, 2007.

[GK97]   P. B. Gibbons and E. Korach. Testing shared memories. *SIAM Journal of Computing*, 26(4):1208–1244, 1997.

[HS12]   Maurice Herlihy and Nir Shavit. *The Art of Multiprocessor Programming*. Morgan Kaufmann, 2012.

[HW90]   M. Herlihy and J. M. Wing. Linearizability: a correctness condition for concurrent objects. *ACM Transactions on Programming Languages and Systems*, 12(3):463–492, 1990.

[Kar72]   Richard M. Karp. Reducibility among combinatorial problems. In Raymond E. Miller, James W. Thatcher, and Jean D. Bohlinger, editors, *Complexity of Computer Computations*, pages 85–103. Springer US, 1972.

[Low16]   Gavin Lowe. Testing for linearizability. *Concurrency and Computation: Practice and Experience*, 29(14), 2016.

[ST05]   Hakan Sundell and Philippas Tsigas. Fast and lock-free concurrent priority queues for multi-thread systems. *Journal of Parallel and Distributed Computing*, 65(5):609–627, 2005.

[WBM+10]   Peter Welch, Neil Brown, James Moores, Kevin Chalmers, and Bernhard Sputh. Alting barriers: synchronisation with choice in Java using JCSP. *Concurrency and Computation: Practice and Experience*, 22(8), 2010.