

Understanding Synchronisation

Jonathan Lawrence and Gavin Lowe

September 27, 2021

Abstract

...

1 Introduction

A common step of many concurrent programs involves two or more threads *synchronising*: each thread waits until other relevant threads have reached the synchronisation point before continuing; in addition, the threads can exchange data. We study synchronisations in this paper.

We start by giving some examples of synchronisations in order to illustrate the idea. (We use Scala notation; we explain non-standard aspects of the language in footnotes.) In each case, the synchronisation is mediated by a *synchronisation object*.

Perhaps the most common form of synchronisation object is a synchronous channel. Such a channel might have signature¹

```
class SyncChan{  
  def send(x: A): Unit  
  def receive(): A  
}
```

Each invocation of one of the operations must synchronise with an invocation of the other operation: the two invocations must overlap in time. If an invocation `send(x)` synchronises with an invocation of `receive`, then the `receive` returns `x`.

For some synchronisation objects, synchronisations might involve more than two threads. For example, an object of the following class

```
class Barrier(n: Int){  
  def sync(): Unit  
}
```

¹The type `Unit` is the type that contains a single value, the *unit value*, denoted `()`.

can be used to synchronise n threads, known as a *barrier synchronisation*: each thread calls `sync`, and no invocation returns until all n have called it.

In addition, we allow the synchronisations to be mediated by an object that maintains some state between synchronisations. As an example, consider a synchronous channel that, in addition, maintains a sequence counter, and such that both invocations receive the value of this counter.

```
class SyncChanCounter{
  private var counter: Int
  def send(x: A): Int
  def receive(): (A, Int)
}
```

We consider what it means for one of these synchronisation objects to be correct. In Section 2 ...

[More here.](#)

2 Specifying synchronisations

In this section we describe how synchronisations can be formally specified. For ease of exposition, we consider just the case of *binary* synchronisation in this section; we generalise in Section 5.

We assume that the synchronisation object has two operations, each of which has a single parameter, with signatures as follows.

```
def op1(x1: A1): B1
def op2(x2: A2): B2
```

(We can model a concrete operation that takes $k > 1$ parameters by an operation that takes a k -tuple as its parameter; we can model a concrete operation that takes no parameters by an operation that takes a `Unit` parameter.) In addition, the synchronisation object might have some state, `state`: S . Each invocation of `op1` must synchronise with an invocation of `op2`, and vice versa. The result of each invocation may depend on the two parameters x_1 and x_2 and the current state. In addition, the state may be updated. The external behaviour is consistent with the synchronisation happening atomically at some point within the duration of both operation invocations (which implies that the invocation must overlap): we refer to this point as the *synchronisation point*.

Each synchronisation object can be specified using a *synchronisation specification object* with the following signature.

```
object Spec{
  def sync(x1: A1, x2: A2): (B1, B2)
```

```
}
```

The idea is that if two invocations $\text{op}_1(x_1)$ and $\text{op}_2(x_2)$ synchronise, then the results y_1 and y_2 of the invocations are such that $\text{sync}(x_1, x_2)$ could return the pair (y_1, y_2) . The specification object might have some private state that is accessed and updated within sync . Note that invocations of sync occur *sequentially*.

We formalise below what it means for a synchronisation object to satisfy the requirements of a synchronisation specification object. But first, we give some examples to illustrate the style of specification.

A typical definition of the specification object might take the following form

```
object Spec{
  private var state: S
  def sync(x1: A1, x2: A2): (B1, B2) = {
    require(guard(x1, x2, state))
    val res1 = f1(x1, x2, state); val res2 = f2(x1, x2, state)
    state = update(x1, x2, state)
    (res1, res2)
  }
}
```

The object has some local state, which persists between invocations. The **require** clause of **sync** specifies a precondition for the synchronisation to take place. The values **res₁** and **res₂** represent the results that should be returned by the corresponding invocations of op_1 and op_2 , respectively.

For example, consider a synchronous channel with operations

```
def send(x: A): Unit
def receive(u: Unit): A
```

(Note that we model the **receive** operation as taking a parameter of type **Unit**, in order to fit our uniform setting.) This can be specified using a synchronisation specification object as follows, with empty state

```
object SyncChanSpec{
  def sync(x: A, u: Unit): (Unit, A) = ((), x)
}
```

If $\text{send}(x)$ synchronises with $\text{receive}()$, then the former receives the unit value $()$, and the latter receives x .

As another example, consider a filtering channel.

```
class FilterChan{
  def send(x: A): Unit
  def receive(p: A => Boolean): A
```

```
}

```

Here the `receive` operation is passed a predicate `p` describing a required property of any value received. This can be specified using a specification object with operation

```
def sync(x: A, p: A => Boolean): (Unit, A) = { require(p(x)); ((), x) }
```

Invocations `send(x)` and `receive(p)` can synchronise only if `p(x)`.

As an example illustrating the use of state in the synchronisation object, recall the synchronous channel with a sequence counter, `SyncChanCounter`, from the introduction. This can be specified using the following specification object.

```
object SyncChanCounterSpec{
  private var counter = 0
  def sync(x: A, u: Unit): (Int, (A, Int)) = {
    counter += 1; (counter, (x, counter))
  }
}
```

2.1 Linearisability

We formalise below precisely the behaviours that should be allowed, given a particular synchronisation specification object. Our definition has much in common with the well known notion of *linearisation* [?], used for specifying concurrent datatypes; so we start by reviewing that notion. There are a number of equivalent ways of defining linearisation: we choose a way that will be convenient subsequently.

A *concurrent history* of an object *o* (either a concurrent datatype or a synchronisation object) records the calls and returns of operation invocations on *o*. It is a sequence of events of the following forms:

- `call.opi(x)`, representing a call of operation *op* with parameter *x*;
- `return.opi:y`, representing a return of an invocation of *op*, giving result *y*.

In each case, *op* is an operation of *o*. Here *i* is a *invocation identity*, used to identify a particular invocation, and to link the `call` and corresponding `return`. In order to be well formed, each invocation identity must appear on at most one `call` event and at most one `return` event; and for each event `return.opi:y`, the history must contain an earlier event `call.opi(x)`, i.e. for the same operation and invocation identity. We consider only well formed histories from now on. We say that a `call` event and a `return` event *match* if they have the same




Figure 1: Timeline representing the linearisation example.

invocation identifier. A concurrent history is *complete* if for every **call** event, there is a matching **return** event, i.e. no invocation is still pending at the end of the history.

For example, consider the following complete concurrent history of a concurrent object that is intended to implement a queue, with operations **enq** and **deq**.

$$h = \langle \text{call.enq}^1(5), \text{call.enq}^2(4), \text{call.deq}^3(), \\ \text{return.enq}^1:(), \text{return.deq}^3:4, \text{return.enq}^2:() \rangle.$$

This history is illustrated by the timeline in Figure 1: here, time runs from left to right; each horizontal line represents an operation invocation, with the left-hand end representing the **call** event, and the right-hand end representing the **return** event.

Linearisability is specified with respect to a specification object *Spec*, with the same operations (and signatures) as the concurrent object in question. A history of the specification object is a sequence of events of the form:

- $op^i(x) : y$ representing an invocation of operation op with parameter x , returning result y ; again i is an invocation identity, which must appear at most once in the history.

A history is *legal* if it is consistent with the definition of *Spec*, i.e. for each invocation, the precondition is satisfied, and the return value is as for the definition of the operation in *Spec*.

For example, consider the history

$$h_s = \langle \text{enq}^2(4):(), \text{enq}^1(5):(), \text{deq}^3:4 \rangle.$$

This is a legal history for a specification object that represents a queue. This history is illustrated by the “ \times ”s in Figure 1.

Let h be a complete concurrent history, and let h_s be a legal history of the specification object corresponding to the same invocations, i.e., for each

$\text{call.op}^i(x)$ and $\text{return.op}^i:y$ in h , h_s contains $\text{op}^i(x):y$, and vice versa. We say that h and h_s are *compatible* if there is some way of interleaving the two histories (i.e. creating a history containing the events of h and h_s , preserving the order of events) such that each $\text{op}^i(x):y$ occurs between $\text{call.op}^i(x)$ and $\text{return.op}^i:y$. Informally, this indicates that the invocations of h appeared to take place in the order described by h_s , and that that order is consistent (in terms of the satisfaction of preconditions and values returned) with the specification object.

Continuing the running example, the histories h and h_s are compatible, as evidenced by the interleaving

$$\langle \text{call.enq}^1(5), \text{call.enq}^2(4), \text{enq}^2(4):(), \text{enq}^1(5):(), \text{call.deq}^3(), \\ \text{return.enq}^1:(), \text{deq}^3:4, \text{return.deq}^3:4, \text{return.enq}^2:() \rangle,$$

which is again illustrated in Figure 1.

We say that a complete history h is *linearisable* with respect to Spec if there is a corresponding valid history h_s of Spec such that h and h_s are compatible.

A concurrent history might not be complete, i.e. it might have some pending invocations. An *extension* of a history h is formed by adding zero or more **return** events corresponding to pending invocations. We write $\text{complete}(h)$ for the subsequence of h formed by removing all **call** events corresponding to pending invocations. We say that a (not necessarily complete) concurrent history h is *linearisable* if there is an extension h' of h such that $\text{complete}(h')$ is linearisable. We say that a concurrent object is linearisable if all of its histories are linearisable.

2.2 Synchronisation linearisability

We now adapt the definition of linearisability to synchronisations. We consider a concurrent history of the synchronisation object *Sync*, as with linearisability; in the case of binary synchronisation, this will contain events corresponding to the operations op_1 and op_2 .

For example, the following is a complete history of the synchronous channel from earlier, and is illustrated in Figure 2:

$$h = \langle \text{call.send}^1(8), \text{call.send}^2(8), \text{call.receive}^3(()), \text{return.receive}^3:(), \\ \text{call.receive}^4(()), \text{return.send}^1:(), \text{call.send}^5(9), \text{return.receive}^4:9, \\ \text{call.receive}^6(()), \text{return.send}^2:(), \text{return.send}^5:(), \text{return.receive}^6:8 \rangle.$$

A history of a synchronisation specification object Spec is a sequence of events of the form




Figure 2: Timeline representing the synchronisation example.

- $\text{sync}^{i_1, i_2}(x_1, x_2) : (y_1, y_2)$, representing an invocation of **sync** with parameters (x_1, x_2) and result (y_1, y_2) . Its invocation identity is (i_1, i_2) : each of i_1 and i_2 must appear at most once in the history.

Such a history is *legal* if is consistent with *Spec*. Informally, a **sync** event with identity (i_1, i_2) represents a synchronisation between the invocations $op_1^{i_1}$ and $op_2^{i_2}$ in a history of the corresponding synchronisation object.

For example, the following is a legal history of **SyncChanSpec**.

$$h_s = \langle \text{sync}^{1,3}(8, ()):(((), 8), \text{sync}^{5,4}(9, ()):(((), 9), \text{sync}^{2,6}(8, ()):(((), 8)) \rangle.$$

The history is illustrated by the “×”s in Figure 2: each event corresponds to the synchronisation of two operations, so is depicted by two (aligned) “×”s on the corresponding operations. This particular synchronisation specification object is stateless, so in fact any permutation of this history would also be legal (but not all such permutations will be compatible with the history of the synchronisation object); but the same will not be true in general of a specification object with state.

Let h be a complete history of the synchronisation object *Sync*. We say that a legal history h_s of *Spec* *corresponds* to h if:

- For each **sync** event with identity (i_1, i_2) in h_s , h contains an invocation of op_1 with identity i_1 and an invocation of op_2 with identity i_2 ;
- For each invocation of op_1 with identity i_1 in h , h_s contains a **sync** event with identity (i_1, i_2) for some i_2 ;
- For each invocation of op_2 with identity i_2 in h , h_s contains a **sync** event with identity (i_1, i_2) for some i_1 .

Informally, a **sync** event with identity (i_1, i_2) represents that the invocations $op_1^{i_1}$ and $op_2^{i_2}$ synchronise.

We say that a complete history h of *Sync* and a corresponding legal history h_s of *Spec* are *synchronisation compatible* if there is some way of interleaving them such that each event $\text{sync}^{i_1, i_2}(x_1, x_2) : (y_1, y_2)$ occurs between $\text{call.op}_1^{i_1}(x_1)$ and $\text{return.op}_1^{i_1} : y_1$, and between $\text{call.op}_2^{i_2}(x_2)$ and $\text{return.op}_2^{i_2} : y_2$.

In the running example, the histories h and h_s are synchronisation compatible, as evidenced by the interleaving illustrated in Figure 2.

We say that a complete history h of *Sync* is *synchronisation linearisable* if there is a corresponding legal history h_s of *Spec* such that h and h_s are synchronisation compatible.

We say that a (not necessarily complete) concurrent history h is *synchronisation linearisable* if there is an extension h' of h such that $\text{complete}(h)$ is synchronisation linearisable. We say that a synchronisation object is synchronisation linearisable if all of its histories are synchronisation linearisable.

Is the definition compositional?

 I think so.

3 Relating synchronisation and linearisation

In this section we describe the relationship between synchronisation linearisation and standard linearisation.

...

It is clear that synchronisation linearisation cannot, in general, be captured directly as standard linearisation. More precisely, given a synchronisation linearisability specification object *SyncSpec*, it is not, in general, possible to find a linearisability synchronisation specification *Spec* such that for every history h , h is synchronisation linearisable with respect to *SyncSpec* if and only if h is linearisable with respect to *Spec*.

For example, consider the example of a synchronous channel from Section 2, where synchronisation linearisation is captured by *SyncChanSpec*. Assume (for a contradiction) that the same property can be captured by linearisation with respect to linearisability specification *Spec*. Consider the history

$$h = \langle \text{call.send}^1(3), \text{call.receive}^2(), \text{return.send}^1():(), \text{return.receive}^2():3 \rangle.$$

This is synchronisation linearisable with respect to *SyncChanSpec*. By the assumption, there must be a legal history h_s of *Spec* such that h and h_s are compatible. Without loss of generality, suppose the **send** in h_s occurs before the **receive**, i.e.

$$h_s = \langle \text{send}^1(3):(), \text{receive}^2():3 \rangle.$$

But the history

$$h' = \langle \text{call.send}^1(3), \text{return.send}^1():(), \text{call.receive}^2(), \text{return.receive}^2():3 \rangle$$

is also compatible with respect to h_s , so h' is linearisable with respect to $Spec$. But then the assumption would imply that h' is synchronisation linearisable with respect to **SyncChanSpec**. This is clearly false, because the operations do not overlap. Hence no such linearisability specification $Spec$ exists.

3.1 Two-step linearisability

In the previous section, we showed that synchronisation linearisation does not correspond directly to linearisation. Nevertheless, we will show that synchronisation linearisability corresponds to a small adaptation of linearisability, but where one of the operations on the concurrent object corresponds to *two* operations of the linearisability specification object. We define what we mean by this, and then prove the correspondence in the next subsection. In the definitions below, we describe just the differences from standard linearisation, to avoid repetition.

Given a synchronisation object with operations **op**₁ and **op**₂, as before, we will consider a linearisability specification object with signature

```
object TwoStepLinSpec{
  def op1(x1: A1): Unit
  def op̄1() : B1
  def op2(x2: A2): B2
}
```

The idea is that the operation **op**₁ on the concurrent object will be linearised by the composition of the two operations **op**₁ and **op̄**₁; but operation **op**₂ on the concurrent object will be linearised by just the operation **op**₂ of the specification object, as before. We call such an object a *two-step linearisability specification object*.

We define a history h_s of such a two-step specification object much as in Section 2.1, except that for each event $\overline{\text{op}}_1^i():y$ in h_s , we require that there is an earlier event $\text{op}_1^i(x):()$ in h_s with the same invocation identity; other than in this regard, invocation identities are not repeated in h_s .

Let h be a complete concurrent history of a synchronisation object, and let h_s be a legal history of a two-step specification object corresponding to the same invocations in the following sense:

- For every **call.op**₁^{*i*}(x) and **return.op**₁^{*i*}: y in h , h_s contains **op**₁^{*i*}(x):() and **op̄**₁^{*i*}(): y ; and vice versa;
- For every **call.op**₂^{*i*}(x) and **return.op**₂^{*i*}: y in h , h_s contains **op**₂^{*i*}(x): y ; and vice versa.

We say that h and h_s are *two-step compatible* if there is some way of interleaving the two histories such that

- Each $\text{op}_1^i(x):()$ and $\overline{\text{op}}_1^i():y$ occur between $\text{call.op}_1^i(x)$ and $\text{return.op}_1^i:y$, in that order;
- Each $\text{op}_2^i(x):y$ occurs between $\text{call.op}_2^i(x)$ and $\text{return.op}_2^i:y$.

For example, consider a synchronous channel, with **send** corresponding to op_1 , and **receive** corresponding to op_2 . Then the following would be an interleaving of two-step compatible histories of the synchronisation object and the corresponding specification object.

$$\langle \text{call.send}^1(3), \text{send}^1(3):(), \text{call.receive}^2(), \text{receive}^2():3, \\ \overline{\text{send}}^1():(), \text{return.send}^1():, \text{return.receive}^2():3 \rangle.$$

The definition of two-step linearisability then follows from this definition of two-step compatibility, precisely as in Section 2.1.


3.2 Proving the relationship

We now prove the relationship between synchronisation linearisation and two-step linearisation.

Consider a synchronisation specification object **SyncSpec**. We build a corresponding two-step linearisation specification object **TwoStepLinSpec** such that synchronisation linearisation with respect to **SyncSpec** is equivalent to two-step linearisation with respect to **TwoStepLinSpec**. The definition is below: the specification's behaviour is described by the automaton on the right.²

```
trait State
case class Zero extends State
case class One(x1: A1) extends State
case class Two(y1: B1) extends State
```

```
object TwoStepLinSpec{
  private var state: State = Zero
  def op1(x1: A1): Unit = {
    require(state.isInstanceOf[Zero]); state = One(x1)
  }
}
```



²Defining the subclasses of **State** as **case classes** allows pattern matching against such values. For example, the statement **val One(x₁) = state** succeeds only if **state** has type **One**, and binds the name x_1 to the value of the x_1 field of **state**.

```

def op2(x2: A2): B2 = {
  require(state.isInstanceOf[One]); val One(x1) = state
  val (y1, y2) = SyncSpec.sync(x1, x2); state = Two(y1); y2
}
def op1(): B1 = {
  require(state.isInstanceOf[Two]); val Two(y1) = state; state = Zero; y1
}
}

```

The definition forces the operations to take place in the order described by the automaton. In addition, the `op2` operation calls the `sync` method on `SyncSpec`, to calculate the return values and to update `SyncSpec`'s state; it stores `op1`'s result in the state.

The following lemma follows immediately from the construction of `TwoStepLinSpec`.

Lemma 1 *Each history of `TwoStepLinSpec` is the concatenation of triples of events of the form `op1i1(x1):()`, `op2i2(x2):y2`, `op1i1():y1` such that `SyncSpec` has a corresponding legal history of events `synci1,i2(x1,x2):(y1,y2)`, and vice versa.*

The following proposition reduces synchronisation linearisability to two-step linearisability.

Proposition 1.1 *Let `SyncObj` be a synchronisation object, `SyncSpec` be a synchronisation specification object, and let `TwoStepLinSpec` be built from `SyncSpec` as above. Then `SyncObj` is two-step linearisable with respect to `TwoStepLinSpec` if and only if it is synchronisation linearisable with respect to `SyncSpec`.*

Proof: (\Rightarrow). Let h be a concurrent history of `SyncObj`. By assumption, there is an extension h' of h , and a legal history h_s of `TwoStepLinSpec` such that $h'' = \text{complete}(h')$ and h_s are two-step compatible. Build a history h'_s of `SyncSpec` by replacing each triple `op1i1(x1):()`, `op2i2(x2):y2`, `op1i1():y1` in h_s by the event `synci1,i2(x1,x2):(y1,y2)`. The history h'_s is legal by Lemma 1. It is possible to interleave h'' and h'_s by placing each event `synci1,i2(x1,x2):(y1,y2)` in the same place as the corresponding event `op2i2(x2):y2` in the interleaving of h'' and h_s ; by construction, this is between `call.op1i1(x1)` and `return.op1i1:y1`, and between `call.op2i2(x2)` and `return.op2i2:y2`. Hence h'' and h_s are synchronisation compatible, so h'' is synchronisation linearisable, and so h is synchronisation linearisable.

(\Leftarrow). Let h be a complete history of `SyncObj`. By assumption, there is an extension h' of h , and a legal history h_s of `SyncSpec` such that $h'' = \text{complete}(h')$ and h_s are synchronisation compatible. Build a history h'_s

of **TwoStepLinSpec** by replacing each event $\text{sync}^{i_1, i_2}(x_1, x_2):(y_1, y_2)$ in h_s by the three events $\text{op}_1^{i_1}(x_1):()$, $\text{op}_2^{i_2}(x_2):y_2$, $\overline{\text{op}}_1^{i_1}():y_1$. The history h'_s is legal by Lemma 1. It is possible to interleave h'' and h'_s by placing each triple $\text{op}_1^{i_1}(x_1):()$, $\text{op}_2^{i_2}(x_2):y_2$, $\overline{\text{op}}_1^{i_1}():y_1$ in the same place as the corresponding event $\text{sync}^{i_1, i_2}(x_1, x_2):(y_1, y_2)$ in the interleaving of h'' and h_s ; by construction, each $\text{op}_1^{i_1}(x_1):()$ and $\overline{\text{op}}_1^{i_1}():y_1$ are between $\text{call.op}_1^{i_1}(x_1)$ and $\text{return.op}_1^{i_1}:y_1$; and each $\text{op}_2^{i_2}(x_2):y_2$ is between $\text{call.op}_2^{i_2}(x_2)$ and $\text{return.op}_2^{i_2}:y_2$. Hence h'' and h_s are two-step compatible, so h'' is two-step linearisable, and so h is two-step linearisable. \square

The two-step linearisation specification object can often be significantly simplified from the template definition above. Here is such a specification object for a synchronous channel.

```
object SyncChanTwoStepLinSpec{
  private var state = 0      // Takes values 0, 1, 2, cyclically
  private var value: A = -   // The current value being sent
  def send(x: A): Unit = { require(state == 0); value = x; state = 1 }
  def receive(u: Unit): A = { require(state == 1); state = 2; value }
  def  $\overline{\text{send}}():$  Unit = { require(state == 2); state = 0 }
}
```

4 Testing algorithms

Overview of linearisability testing framework.

Note, assumes deterministic specification object.

Desirable outcome: no false errors; real errors detected with high probability (given enough tests). Distinguish between the real history, in terms of calls and returns from operations, and the corresponding log, which might contain delays.

4.1 Hacking the linearisability framework

Can we encode synchronisation linearisability within the linearisability tester, using the correspondence of the previous section? Jonathan to think about this. Will need to perform two log additions for some concrete operations.


The idea of linearisability testing is as follows. We run several threads, performing operations (typically chosen randomly) upon the concurrent datatype that we are testing, and logging the calls and returns. More precisely, a thread that performs a particular operation $\text{op}(x)$: (1) writes $\text{call.op}(x)$ into the log; (2) performs $\text{op}(x)$ on the synchronisation object, obtaining result y , say; (3) writes $\text{return.op}:y$ into the log. Once all threads have

finished, we can use an algorithm to test whether the history is linearisable. This can be repeated many times.

Note that there is potentially a delay between writing the **call** event into the log and actually calling the operation; and likewise there is potentially a delay between the operation returning and writing the **return** event into the log. However, these delays do not generate false errors: if a history without such delays is linearisable, then so is a corresponding history with delays. Our experience is that the delays do not prevent the detection of bugs when they exist (although might require performing the test more times).

It turns out that we cannot use linearisability testing directly with the specification object **TwoLinSpec** from Section 3.2. The obvious way to try to do this would be as follows. A thread that performs the concrete operation $\text{op}_1(x_1)$: (1) writes $\text{call.op}_1^i(x_1)$ into the log; (2) performs $\text{op}_1(x_1)$ on the synchronisation object, obtaining result y_1 ; (3) writes $\text{return.op}_1^i()$, $\text{call.op}_1^i()$ and $\text{return.op}_1^i:y_1$ into the log. A thread that performs operation op_2 acts as for standard linearisation testing. Once all threads have finished, we can use the existing algorithms for testing whether the history is linearisable with respect to **TwoStepLinSpec**.

This approach does not work, because it gives false errors. For example, the timeline below depicts a log that could be obtained from a correct synchronous channel using the above approach, where we treat **send** as op_1 .



Here, the top two threads synchronise to transmit 5, and then the bottom two threads synchronise to transmit 6. However, the top thread is slow to write its last three events into the log. The above history is not linearisable with respect to **TwoStepLinSpec**: this would require send^1 to be linearised before $\text{send}^3(6)$. Hence the approach would generate a false error.

Instead we use a technique that is robust against delays in logging. We assume that each thread has an identity in some range $[0 \dots \text{NumThreads})$. We arrange for this identity to be included in the **call** events written to the log for operations op_1 and $\overline{\text{op}}_1$, but otherwise threads act as above. In particular, this means that for each thread, calls to op_1 and $\overline{\text{op}}_1$ alternate.

We then test whether the history is linearisable with respect to the specification object below. It requires that corresponding invocations of op_1 and op_2 are linearised consecutively, but allows the corresponding $\overline{\text{op}}_1$ to be linearised later (but before the next operation invocation by the same thread). It uses an array `returns`, indexed by thread identities, to record values that should be returned by a $\overline{\text{op}}_1$ operation.

```

type ThreadID = Int           // Thread identifiers
val NumThreads: ThreadID = ... // Number of threads
trait State
case class Zero extends State
case class One(t: ThreadID, x1: A1) extends State

object TwoStepDelayedLinSpec{
  private var state: State = Zero
  private val returns = new Array[B1](NumThreads)
  def op1(t: ThreadID, x1: A1): Unit = {
    require(state.isInstanceOf[Zero]); state = One(t, x1)
  }
  def op2(x2: A2): B2 = {
    require(state.isInstanceOf[One]); val One(t, x1) = state
    val (y1, y2) = SyncSpec.sync(x1, x2); returns(t) = y1; state = Zero
  }
  def  $\overline{\text{op}}_1$ (t: ThreadID): B1 = returns(t)
}

```

In order to argue for correctness, we need to distinguish between:

- the *invocation history*, in terms of actual calls and returns of op_1 and op_2 ; and
- the corresponding *log history*, which might contain delays.

Further, the invocation history uses events of the form $\text{call.op}_1^{i_1}(x_1)$, $\text{return.op}_1^{i_1}:y_1$, $\text{call.op}_2^{i_2}(x_2)$, and $\text{return.op}_2^{i_2}:y_2$; and the log history uses events of the form $\text{call.op}_1^{i_1}(t, x_1)$ (note the additional thread identity parameter), $\text{return.op}_1^{i_1}():()$ (note the unit return value), $\text{call}.\overline{\text{op}}_1^{i_1}(t)$, $\text{return}.\overline{\text{op}}_1^{i_1}:y_1$ (note the transferred return value), $\text{call.op}_2^{i_2}(x_2)$, and $\text{return.op}_2^{i_2}:y_2$.

We can consider the interleaving of the histories, following real-time order. In the interleaving, for op_1 and op_2 , **call** events from the log history may be earlier than the corresponding **call** events from the invocation history; and **return** events from the log history may be later than the corresponding **return** events from the invocation history. However, the two histories agree on the relative order of the op_1 and op_2 events of each individual thread.

The lemma below shows that this testing method does not generate any false errors.

Lemma 2 *Suppose a complete invocation history h is synchronisation linearisable with respect to **syncSpec**. Then each corresponding log history h_l is linearisable with respect to **TwoStepDelayedLinSpec**.*

Proof: ** Non-complete history

Consider the invocation history h and a corresponding log history h_l ; and consider their real-time interleaving. By assumption, there is a legal history h_s of **SyncSpec** such that h and h_s are synchronisation compatible. Thus h_s may be interleaved with the interleaving of h and h_l , so that each **sync** event from h_s is placed between the corresponding **call** and **return** events from h , and hence also between the corresponding **call** and **return** events from h_l .

We build an interleaving of h , h_l and a legal history h'_s of **TwoStepDelayedLinSpec** from the interleaving of h , h_l and h_s , as follows.

1. First, we replace each **sync** ^{i_1, i_2} (x_1, x_2):(y_1, y_2) by the two (consecutive) events **op**₁ ^{i_1} (t, x_1):() and **op**₂ ^{i_2} (x_2): y_2 , where t is the identity of the thread that makes the corresponding call of **op**₁ in h_l .
2. Secondly, we insert an event **op**₁ ^{i_1} (t): y_1 between every **call.op**₁ ^{i_1} (t) and **return.op**₁ ^{i_1} : y_1 (from h_l), but not between any pair of **op**₁ and **op**₂ events from the previous stage.

Note that each inserted event is between the corresponding **call** and **return** events from h_l , by construction. Let h'_s be these inserted events; we show that h'_s is a legal history of **TwoStepDelayedLinSpec**.

- The events inserted in step 1 alternate between **op**₁ and **op**₂, as required by **TwoStepDelayedLinSpec**. Further, they are in the same order, and have the same value for y_2 , as the corresponding **sync** events from h_s . Hence each inserted **op**₂ event has the return value y_2 as required by **TwoStepDelayedLinSpec**. Finally the value y_1 written into **returns**(t) (by **op**₂ in **TwoStepDelayedLinSpec**) matches the value returned by the corresponding call to **sync**.
- For the events from step 2, because of the way logging is done, each value of y_1 returned by **op**₁ must match the value returned by the previous invocation of **op**₁ on the synchronisation object by thread t . Since h is synchronisation linearisable, this y_1 must match the value returned by the corresponding call of **sync**. And this matches the last value written into **returns**(t) (by the previous item), as required by **TwoStepDelayedLinSpec**.

This demonstrates that h_l is linearisable with respect to **TwoStepDelayedLinSpec**. \square

In theory, the delays in logging may mean that an invocation history that is not synchronisation linearisable is transformed into a log history that is linearisable with respect to **TwoStepDelayedLinSpec** (although this seems unlikely). However, if the delays are sufficiently small, the log history agrees with the invocation history on the **op**₁ and **op**₂ events. We show that in this case that log history is not linearisable.

Lemma 3 *Consider an invocation history h that is not synchronisation linearisable with respect to **SyncSpec**. Let h_l be a corresponding log history that agrees with h on **op**₁ and **op**₂ events. Then h_l is not linearisable with respect to **TwoStepDelayedLinSpec**.*

Proof: We prove the contrapositive: we suppose that h_l is linearisable with respect to **TwoStepDelayedLinSpec**, and show that h is synchronisation linearisable with respect to **SyncSpec**. \square

4.2 Case with state

Suppose the specification object has non-trivial state.

I think it will be more efficient to give a more direct implementation. Define a configuration to be: (1) a point in the log reached so far; (2) the set of pending operation invocations that have not synchronised; (3) the set of pending operation invocations that have synchronised (but not returned); and (4) the state of the sequential synchronisation object. In any configuration, can: synchronise a pair of pending operations (and update the synchronisation object); advance in the log if the next event is a return that is not pending; or advance in the log if the next event is a call. Then perform DFS.

Partial order reduction: a synchronisation point must follow either the call of one of the concurrent operations, or another synchronisation point. Any synchronisation history can be transformed into this form, by moving synchronisation points earlier, but not before any of the corresponding call events, and preserving the order of synchronisations. This means that after advancing past the call of an invocation, we may synchronise that invocation, and then an arbitrary sequence of other invocations.

Alternatively, a synchronisation point must precede either the return of one of the concurrent operations, or another synchronisation point. This is more like the JIT technique in the linearisability testing paper. This means

that before advancing in the log to the return of an invocation that has not synchronised, we synchronise some invocations, ending with the one in question. And we only synchronise in these circumstances.

My intuition is that the former is more efficient: in the latter, we might investigate synchronising other invocations even though the returning operation can't be synchronised with any invocation.

Complexity

Consider the problem of testing whether a given concurrent history has synchronisations consistent with a given sequential specification object.

We make use of a result from [?] concerning the complexity of the corresponding problem for linearizability. Let **Variable** be a linearizability specification object corresponding to a variable with **get** and **set** operations. Then the problem of deciding whether a given concurrent history is linearisable with respect to **Variable** is NP-complete.

Let **ConcVariable** be a concurrent object that represents a variable.

We consider concurrent synchronisation histories on an object with the following signature.

```
object VariableSync{
  def op1(op: String, x: Int): Int
  def op2(u: Unit): Unit
}
```

The intention is that **op₁("get", x)** acts like **get(x)**, and **op₁("set", x)** acts like **set(x)** (but returns -1). The **op₂** invocations do nothing except synchronise. This can be captured formally by the following synchronisation specification object.

```
object VariableSyncSpec{
  private var state = 0
  def sync((op, x): (String, Int), u: Unit): (Int, Unit) =
    if (op == "get") (state, ()) else { state = x; (-1, ()) }
}
```

Let **ConcVariable** be a concurrent object that represents a variable. Given a concurrent history h of **ConcVariable**, we build a concurrent history h' of **VariableSync** as follows. We replace every call or return of **get(x)** by (respectively) a call or return of **op₁("get", x)**; and we do similarly with **sets**. If there are k calls of **get** or **set** in total, we prepend k calls of **op₂**, and append k corresponding returns (in any order). Then it is clear that h is linearisable with respect to **Variable** if and only if h' is linearisable with respect to **VariableSyncSpec**.

4.3 Stateless case

In the stateless case, a completely different algorithm is possible. Define two invocations to be compatible if they could be synchronised, i.e. they overlap and the return values agree with those for the specification object. For n invocations of each operation (so a history of length $4n$), this can be calculated in $O(n^2)$. Then find if there is a total matching in the corresponding bipartite graph, using the Ford-Fulkerson method, which is $O(n^2)$.

5 Variations

We've implicitly assumed that the operations op_1 and op_2 are distinct. I don't think there's any need for this. Example: exchanger.

Most definitions and results go through to the case of $k > 2$ invocations synchronising. Examples: ABC problem; barrier synchronisation. To capture the relationship with linearisation, we require $k - 1$ operations to be linearised by two operations of the specification object. Maybe give automaton for $k = 4$.

It turns out that for $k > 2$, the problem of deciding whether a history is synchronisation linearisable is NP-complete in general, even in the stateless case. We prove this fact by reduction from the following problem, which is known to be NP-complete ??.

Definition 3.1 *The problem of finding a complete matching in a 3-partite hypergraph is as follows: given finite sets X , Y and Z of the same cardinality, and a set $T \subseteq X \times Y \times Z$, find $U \subseteq T$ such that each member of X , Y and Z is included in precisely one element of U .*

Suppose we are given an instance (X, Y, Z, T) of the above problem. We construct a synchronisation specification and a corresponding history h such that h is synchronisation linearisable if and only if a complete matching exists. The synchronisations are between operations as follows:

```
def op1(x: X): Unit
def op2(y: Y): Unit
def op3(z: Z): Unit
```

The synchronisations are specified by:

```
def sync(x: X, y: Y, z: Z): (Unit, Unit, Unit) = {
  require((x, y, z) ∈ T); ((), (), ())
}
```

The history h starts with calls of $\text{op}_1(x)$ for each $x \in X$, $\text{op}_2(y)$ for each $y \in Y$, and $\text{op}_3(z)$ for each $z \in Z$ (in any order); and then continues with returns of the same invocations (in any order). It is clear that any synchronisation linearisation corresponds to a complete matching, i.e. the invocations that synchronise correspond to the complete matching U .

5.1 Different modes of synchronisation

Some synchronisation objects allow different modes of synchronisation. For example, consider a synchronous channel with timeouts: each invocation might synchronise with another invocation, or might timeout without synchronisation. Such a channel might have a signature as follows.

```
class TimeoutChannel{
  def send(x: A): Boolean
  def receive(u: Unit): Option[A]
}
```

The `send` operation returns a boolean to indicate whether the send was successful, i.e. whether it synchronised. The `receive` operation can return a value `Some(x)` to indicate that it synchronised and received `x`, or can return the value `None` to indicate that it failed to synchronise (the type `Some[A]` contains the union of such values). The possible synchronisations can be captured by the following specification object.

```
object TimeoutSpec{
  def syncs,r(x: A, u: Unit): (Boolean, Option[A]) = (true, Some(x))
  def syncs(x: A): Boolean = false
  def syncr(u: Unit): Option[A] = None
}
```

The operation `syncs,r` corresponds to where a `send` and `receive` synchronise, as previously. The operations `syncs` and `syncr` correspond, respectively, to where a `send` or `receive` fails to synchronise.

More generally, the specification object can have any number of operations of the form

```
def syncj1,...,jm(x1: A1, ..., xm: Am): (B1, ..., Bm)
```

This corresponds to the case of a synchronisation between the m invocations $\text{op}_{j_1}(x_1), \dots, \text{op}_{j_m}(x_m)$. The formal definition is an obvious adaptation of the previous version: in the interleaved history, between the call and return of each $\text{op}_j(x) : y$, there must be a corresponding `syncj1,...,jm(x1, ..., xm) : (y1, ..., ym)` event, i.e. for some i , $j = j_i$, $x = x_i$, and $y = y_i$.

*** Can we capture the bounded buffer example in this framework?